

3 Compilers

Interpreters and compilers:

- * Assume a high-level program P takes input I and produces output O .
- * An interpreter takes P and I and produces O .
- * Compiling breaks this process into two parts.
- * The compiler takes P and produces low-level machine code M .
- * The processor takes M and I and produces O .



Remarks:

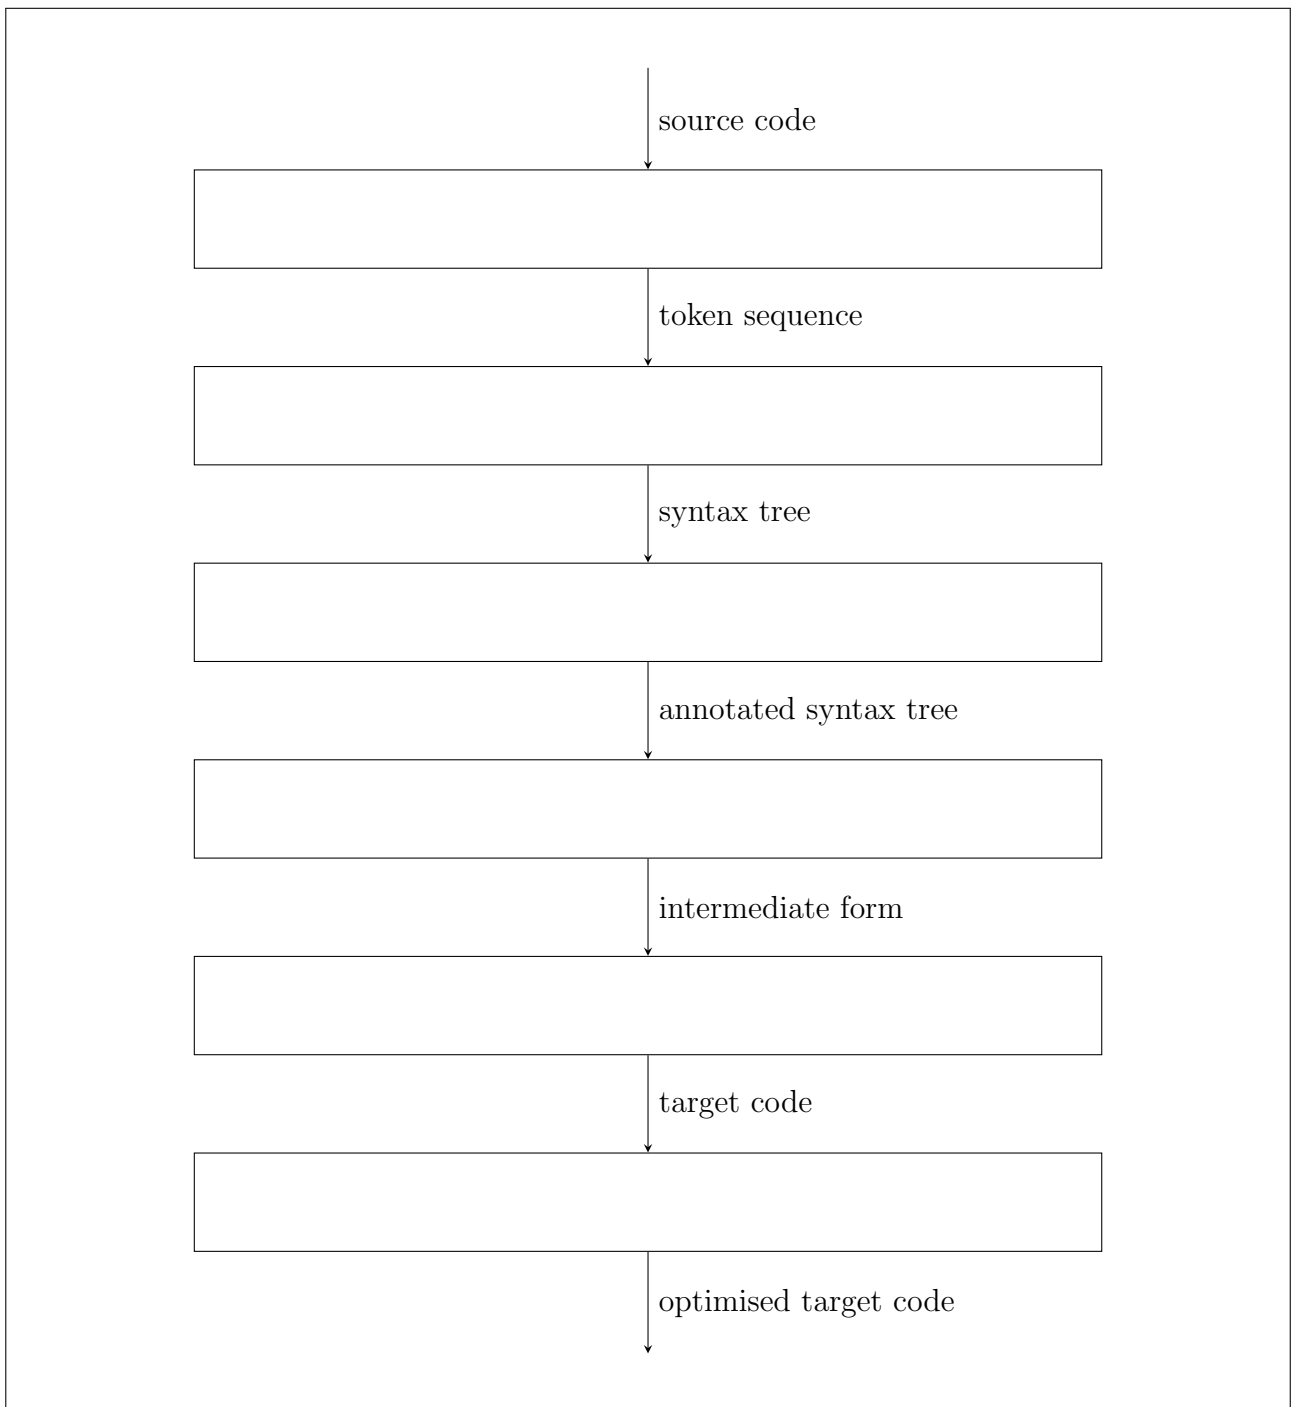
- * The processor is an interpreter for a low-level language.
- * The processor might be a virtual machine; see Java.
- * Libraries may be part of the run-time environment or linked after compiling.

Benefits of compiling:

- * The compiler performs tasks once, which would have to be repeated during interpretation.
- * Analyse programs to check aspects of correctness.
- * Optimise programs so they use less time and space.

Structure of a compiler:

- * The input of a compiler is a sequence of characters: the source-level program.
- * The *syntax tree* constructed from this sequence reflects the structure of the program.
- * Analyses such as type checking and some optimisations can be performed on this tree.
- * The tree is ultimately converted to a low-level instruction sequence such as assembly code.
- * A compiler is typically divided into phases.



Topics discussed in the following:

- * lexical analysis
- * syntax analysis
- * semantic analysis
- * optimisation
- * code generation for virtual machines

Topics not discussed are:

- * code generation for real machines
- * error-handling

The diagram below exemplifies the analysis:

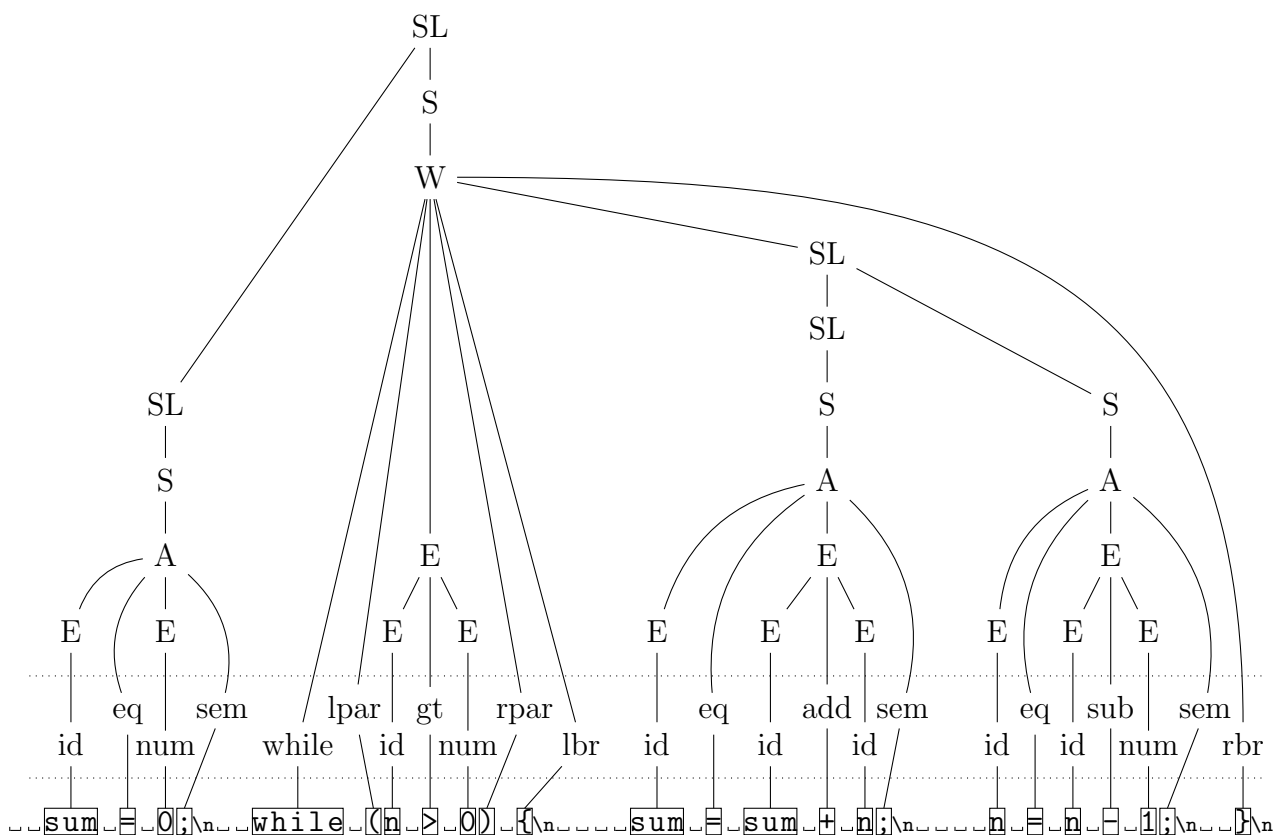
* Consider the following fragment of a C program or a Java program:

```
sum = 0;
while (n > 0) {
    sum = sum + n;
    n = n - 1;
}
```

* At the bottom of the diagram is the sequence of characters of this program.

* Above that is the sequence of *tokens* generated by the scanner.

* On top is the syntax tree generated by the parser.



The nodes of the syntax tree have the following meaning:

- * SL: statement list
- * S: statement
- * W: while-loop
- * A: assignment
- * E: expression
- * id, eq, num, sem, while, lpar, gt, rpar, lbr, add, sub, rbr: tokens of the source program
- * Tokens can be annotated with their position in the source code.
- * White-space characters are eliminated by the scanner.

The syntax tree is translated to code for a low-level machine:

- * The machine holds the values of variables in a small number of *registers*.
- * Calculations take place on a *stack*.
- * The contents of registers are pushed to and popped from the stack.
- * The processor executes the code step-by-step unless instructed to jump:

```
        sipush 0      -- constant 0
        istore 0      -- register 0 holds the value of sum
11:     iload 1        -- register 1 holds the value of n
        sipush 0
        if_icmple 12  -- jump to 12 if n <= 0
        iload 0
        iload 1
        iadd
        istore 0      -- completes sum = sum + n
        iload 1
        sipush 1
        isub
        istore 1      -- completes n = n - 1
        goto 11
12:
```

3.1 Lexical Analysis

The first step of compiling is performed by the *scanner*:

- * It reads the source code character-by-character.
- * It produces a sequence of *tokens* to be fed to the parser.
- * White-space and comments are typically discarded.
- * It keeps track of source-code line and column for layout checking and error messages.

There are several kinds of token:

- * reserved words, such as
- * symbols, such as
- * identifiers, such as variable names
- * number constants, such as
- * string constants, such as
- * Identifiers and constants have the concrete value attached.

Use of the scanner:

- * The scanner will not convert the entire source code at once.
- * Typically, it is called by the parser to deliver the next token in the source.
- * The parser needs *lookahead* that shows the next token without consuming it.

Each kind of token is described by a regular expression:

- * For reserved words and symbols, this might be just a simple string.
- * For other kinds, a sequence of definitions might be used:

$$\begin{aligned} digit &= 0-9 \\ lower &= a-z \\ upper &= A-Z \\ letter &= lower|upper \\ identifier &= \\ nonzerodigit &= 1-9 \\ decinteger &= \\ hexdigit &= digit|a-f|A-F \\ hexinteger &= \\ integer &= decinteger|hexinteger \\ stringchar &= \text{any character except " and \ and newline} \\ escape &= \backslash(\text{any character}) \\ stringitem &= stringchar|escape \\ string &= \end{aligned}$$

- * Such a sequence of definitions must not be cyclic.
- * It can be converted to a flat regular expression by repeated substitution.
- * *Character classes* abbreviate choices: $a-z$ amounts to $a|b|\dots|y|z$.
- * The resulting regular expression is converted to a DFA that is used for matching.

3.1.1 Scanners and Automata

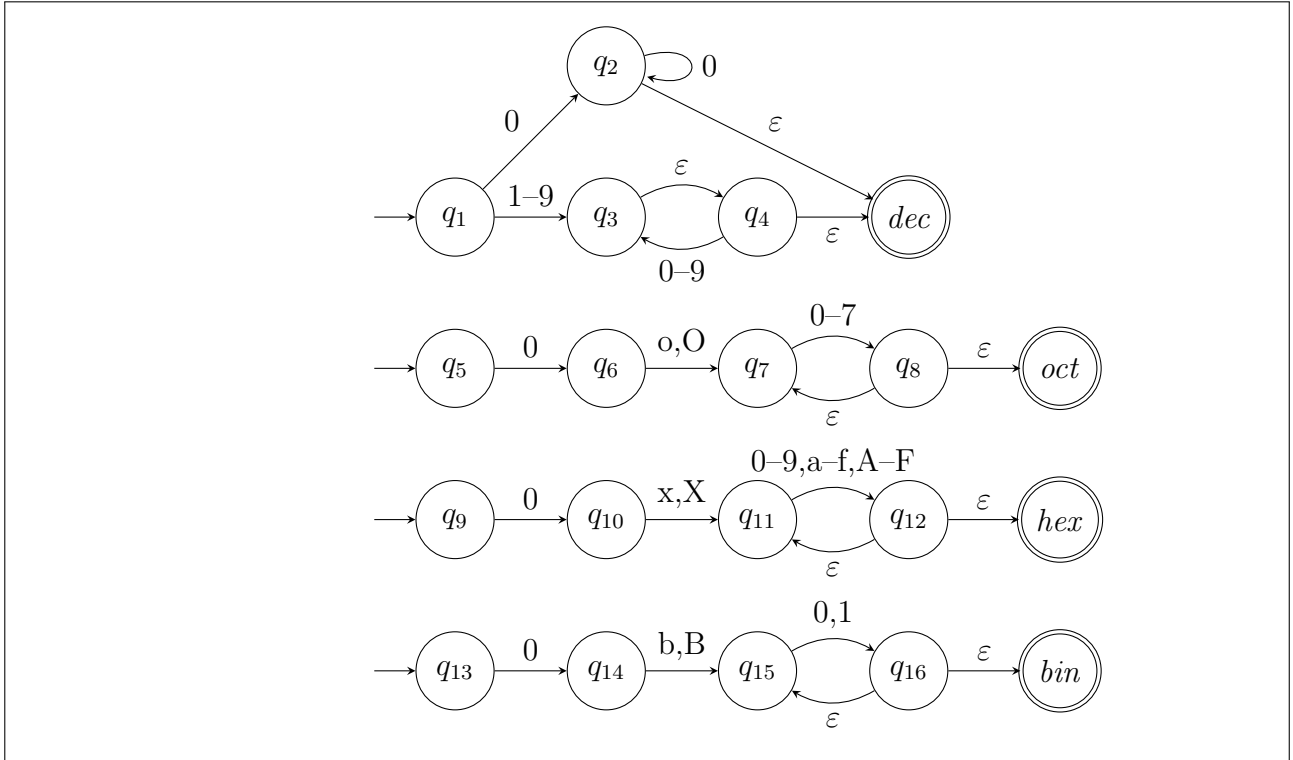
The scanner recognises several kinds of token at the same time.

- * For example, there might be four tokens for integers in different bases:

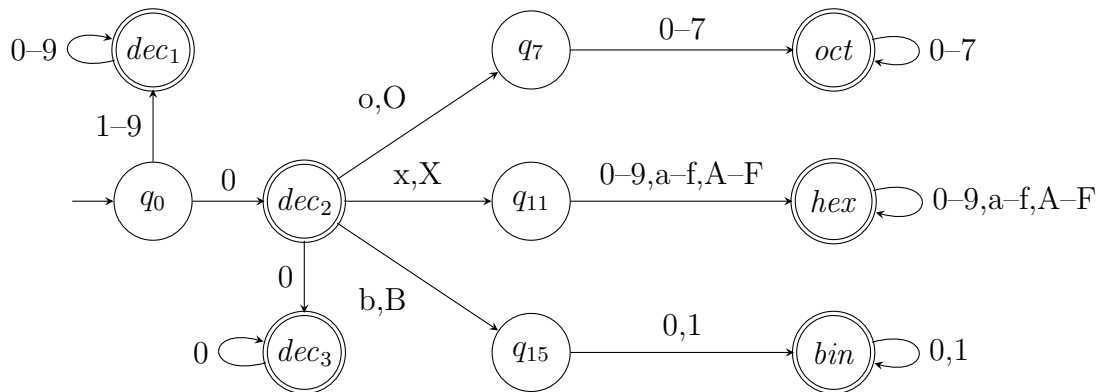
$$\begin{aligned} decinteger &= 0^+|nonzerodigit\ digit^* \\ octinteger &= 0(o|O)octdigit^+ \\ hexinteger &= 0(x|X)hexdigit^+ \\ bininteger &= 0(b|B)bindigit^+ \\ octdigit &= 0-7 \\ bindigit &= 0|1 \end{aligned}$$

- * An NFA is constructed for each kind of token that needs to be recognised.
- * Each accept state is marked with the respective token.

- * A new initial node is connected by ε -transitions to the initial node of each NFA.



- * The resulting NFA is converted to a DFA and minimised.



- * Marks on the accept states are preserved by the subset construction.
- * They identify which token has been recognised.

Operation of the scanner:

- * In each step, the scanner tries to match as much of the input as possible.
- * The DFA consumes input as long as there is more input and a transition is possible.
- * If the resulting state is accepting, the corresponding token is returned.
- * The user specifies which token to return if more than one corresponds to a state.
- * If the resulting state is not accepting, the DFA returns to the last accept state.
- * To this end, transitions are undone and corresponding symbols are put back to the input.
- * If no accept state is reached, the scanner reports a lexical error.
- * The DFA is repeatedly restarted on the remaining input to find further tokens.

3.1.2 Extensions

Extended regular expressions:

- * Eliminate at the regular expression stage:
- * Reducing p^+ to pp^* can cause
- * Special automata may be devised in such cases; for p^+ construct

- * Eliminate at the automaton stage: for \bar{p} ,
- * Some constructs require more complex constructions.

Consider C-style comments `/* ... */`:

- * Unlike typical strings, they are delimited by multiple characters.
- * The characters in the comment cannot be described using character classes.
- * A regular expression for such comments is $/*\overline{\Sigma^*}*/$.

Constructing a minimal DFA for

The above procedure can be automated.

- * Describe lexical structure by a sequence of definitions using extended regular expressions.
- * A *scanner generator* takes such a description and produces a scanner.
- * Examples of scanner generators are Lex, Flex, JFlex, PLY.
- * Each kind of token may have an action which is performed when the token is found.

3.2 Syntax Analysis

The second step of compiling is performed by the *parser*:

- * It calls the scanner to deliver tokens as required.
- * It produces a *syntax tree* for further analysis and code generation.
- * Lookahead tokens are used to efficiently recognise the structure.
- * The output is typically an *abstract syntax tree* that omits irrelevant details.
- * The parser may generate code on-the-fly, without constructing the full syntax tree.

The syntactic structure of a program is described in *Backus-Naur form* (BNF):

- * A BNF description is a set of grammar rules.
- * Each rule describes the structure of a program fragment.
- * Unlike sequences of definitions using regular expressions, the rules may be *recursive*.
- * Examples of fragments are statements, expressions, declarations, parameter lists.
- * Rules for expressions, comparisons and statements might read as follows:

Expression =

Arithmetic =

Comparison = *Expression Relation Expression*

Relation = = | != | < | <= | > | >=

Statements =

Statement = *If* | *While* | *Assignment*

Assignment =

While = **while** *Comparison* **do** *Statements* **end**

If =

- * An item on the left-hand side of a rule is a *non-terminal*.
- * It may occur any number of times on the right-hand side of any rule.
- * Several rules for the same non-terminal are abbreviated by | on the right-hand side:

Expression = *Expression Arithmetic Expression*

Expression = (*Expression*)

Expression = **number**

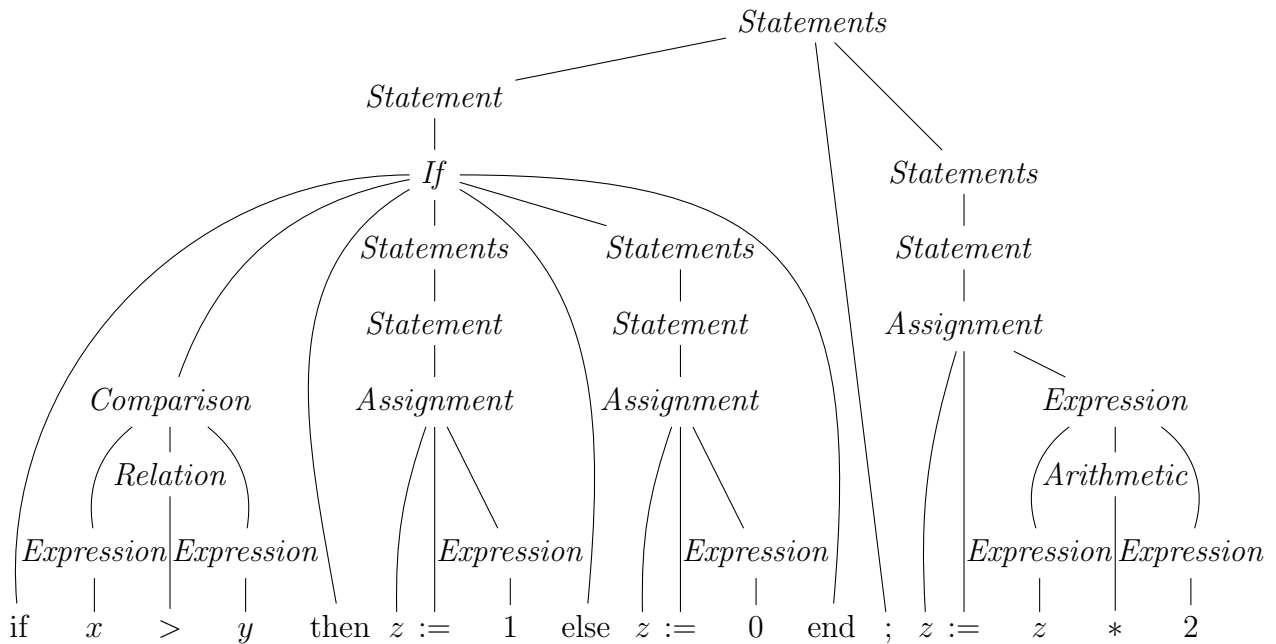
Expression = **identifier**

- * This represents a choice, as for regular expressions.
- * Any other item on the right-hand side is a *terminal*.
- * Terminals are matched by the tokens delivered from the scanner.

3.2.1 Syntax Trees


The parser constructs a syntax tree from a sequence of tokens.

- * Every leaf of the syntax tree is labelled with a terminal.
- * Every inner node of the syntax tree corresponds to the application of a rule.
- * The inner node is labelled with the non-terminal on the left-hand side.
- * Its children are labelled with the items on the right-hand side.
- * The tree is *ordered*: the sequence of children matters.
- * An example of a syntax tree is:



Consider the input `2 + 8 * 5`.

- * The scanner delivers the token sequence `number(2), +, number(8), *, number(5)`.
- * A syntax tree for this token sequence is:




- * Another syntax tree for the same token sequence is:



- * The BNF grammar is *ambiguous* for the given token sequence.
- * Evaluation of the expression gives different values depending on the syntax tree.
- * The programmer might use parentheses to specify precedence.
- * The language might specify precedence rules that need to be reflected in the BNF:

$$\begin{aligned}
 \textit{Expression} &= \textit{Term} \mid \textit{Expression Additive Term} \\
 \textit{Additive} &= + \mid - \\
 \textit{Term} &= \textit{Factor} \mid \textit{Term Multiplicative Factor} \\
 \textit{Multiplicative} &= * \mid / \\
 \textit{Factor} &= (\textit{Expression}) \mid \textbf{number} \mid \textbf{identifier}
 \end{aligned}$$

- * With this BNF the token sequence has just one syntax tree:



- * Intermediate nodes, which help parsing, might be omitted in an abstract syntax tree.

Parsers construct the syntax tree

- * top-down: LL, recursive-descent;
- * bottom-up: LR, LALR, SLR, shift-reduce.
- * LR is more expressive than LL by deferring decisions until more information is available.
- * Recursive-descent parsers are easier to understand than shift-reduce parsers.

3.2.2 Extensions

Extended BNF has regular expressions on the right-hand side.

- * The rules for expressions and terms can be written as:

$$\begin{aligned} \textit{Expression} &= (\varepsilon \mid \textit{Expression Additive}) \textit{Term} \\ \textit{Term} &= (\varepsilon \mid \textit{Term Multiplicative}) \textit{Factor} \end{aligned}$$

- * Using $[p]$ for optional p , this is:

$$\begin{aligned} \textit{Expression} &= [\textit{Expression Additive}] \textit{Term} \\ \textit{Term} &= [\textit{Term Multiplicative}] \textit{Factor} \end{aligned}$$

- * An alternative extended BNF is:

$$\begin{aligned} \textit{Expression} &= \textit{Term} (\textit{Additive Term})^* \\ \textit{Term} &= \textit{Factor} (\textit{Multiplicative Factor})^* \end{aligned}$$

Syntax diagrams are a graphical representation of extended BNF:

- * The rule $\textit{Factor} = (\textit{Expression}) \mid \texttt{number} \mid \texttt{identifier}$ is represented as:

- * The rule $\textit{Term} = [\textit{Term Multiplicative}] \textit{Factor}$ is represented as:

* The rule $Expression = Term (Additive\ Term)^*$ is represented as:

* Rectangles represent non-terminals and ovals represent terminals.

* Arrows represent sequence, choice and iteration.

3.2.3 Recursive-Descent Parsers

A parser can be implemented as a set of mutually recursive functions.

* Consider the extended BNF for expressions:

$$\begin{aligned} Expression &= Term ((+ | -) Term)^* \\ Term &= Factor ((* | /) Factor)^* \\ Factor &= (Expression) | \text{number} | \text{identifier} \end{aligned}$$

* It results in the following parser:

```
def expression():
    term()
    while lookahead() in [ADD, SUB]:
        consume(ADD, SUB)
        term()

def term():
    factor()
    while lookahead() in [MUL, DIV]:
        consume(MUL, DIV)
        factor()

def factor():
    if lookahead() == LPAR:
        consume(LPAR)
        expression()
        consume(RPAR)
    elif lookahead() == NUM:
        consume(NUM)
    elif lookahead() == ID:
        consume(ID)
    else:
        raise Exception
```

The parser calls the scanner to obtain tokens.

* ADD, SUB, MUL, DIV, LPAR, RPAR, NUM, ID are tokens.

* Two functions form the interface to the scanner:

```
def lookahead():
    '''Returns the next token without consuming it.'''

def consume(*expected_tokens):
    '''Consumes the next token, if it is in expected_tokens.
       Raises an exception otherwise.'''
```

Extended BNF can be translated to a *recursive-descent* parser:

* Every non-terminal n is translated to a function.

* The function's body is obtained from the right-hand side r of the BNF rule $n = r$.

```
def n():
    parse(r)
```


Let $parse(r)$ be the parser code for regular expression r .

* $parse(r)$ is defined by induction over the structure of r .

* $parse(q|r)$ is the conditional

* $parse(r^*)$ is the while-loop

* $parse(qr)$ is the sequence



* If r is a non-terminal, then $parse(r)$ is the function call

$r()$

* If r is a terminal, then $parse(r)$ is a call to the scanner to consume token r

`consume(r)`

The scanner needs to look ahead to inform the decisions in recursive-descent parsing:

- * The scanner might provide a lookahead of one token.
- * It is checked whether a string described by the regular expression begins with that token.
- * The *First* set of r comprises the symbols that can begin a string described by r .
- * If $\varepsilon \in L(r)$, then any symbol possibly following r must be considered.
- * The *Follow* set of an occurrence of r is the symbols that can follow a string described by r .
- * There are algorithms to compute the First and Follow sets from an extended BNF.

3.2.4 Abstract Syntax Trees

An abstract syntax tree can be constructed while parsing takes place.

* The interface to the scanner is extended thus:

```
def consume(*expected_tokens):  
    '''Returns the next token and consumes it, if it is in  
    expected_tokens. Raises an exception otherwise.  
    If the token is a number or an identifier, not just the  
    token but a pair of the token and its value is returned.'''
```

* Expressions can be represented by the following classes:

```
class Expression_AST:  
    def __init__(self, left, op, right):  
        self.left = left  
        self.op = op  
        self.right = right  
  
class Number_AST:  
    def __init__(self, number):  
        self.number = number  
  
class Identifier_AST:  
    def __init__(self, identifier):  
        self.identifier = identifier
```

* The parser constructs and returns objects of the above classes:

```
operator = { ADD: '+', SUB: '-', MUL: '*', DIV: '/' }

def expression():
    result = term()
    while lookahead() in [ADD, SUB]:
        op = consume(ADD, SUB)
        tree = term()
        result = Expression_AST(result, operator[op], tree)
    return result

def term():
    result = factor()
    while lookahead() in [MUL, DIV]:
        op = consume(MUL, DIV)
        tree = factor()
        result = Expression_AST(result, operator[op], tree)
    return result

def factor():
    if lookahead() == LPAR:
        consume(LPAR)
        result = expression()
        consume(RPAR)
        return result
    elif lookahead() == NUM:
        value = consume(NUM)[1]
        return Number_AST(value)
    elif lookahead() == ID:
        value = consume(ID)[1]
        return Identifier_AST(value)
    else:
        raise Exception
```

* There is just one kind of arithmetic expression node in the abstract syntax tree.

* Irrelevant details, such as parenthesis tokens, are omitted from the tree.

3.3 Semantic Analysis

The third step of compiling is performed by the *semantic analyser*.

* It traverses the abstract syntax tree constructed by the parser.

* Nodes are annotated with information about the program.

* This information can be used for optimisation, code generation and error handling.

* The *semantics* of a program is its meaning.

* Semantic analysis covers just a few aspects of the meaning.

* *Static semantics* are those aspects that can be checked at compile-time.

Examples of semantic analyses are:

- * type checking: are functions applied to arguments of matching type?
- * type inference: what is the type of an expression?
- * declaration: is every variable declared exactly once?
- * definite assignment: is a variable assigned before it is used?
- * binding: to which declaration does the use of a variable belong?

3.3.1 Type Systems

Many functions are partial, that is, they cannot be applied to all arguments:

- * Add two numbers, but not two lists.
- * Access the first element of a list, but not of a set.
- * Compare two integers for \leq , but not two complex numbers.
- * Find the shortest path in a graph, but not in a string.
- * Sort a list of integers, but not a list that mixes integers and strings.
- * Convert an integer to a string, but not a function.

The *type* of a variable or expression is the set of possible values it can take.

- * A function might declare the types of its parameters and the type of its result.
- * Such a function can only be applied to expressions with matching type.
- * Its result can only be used in a context with matching type.

Low-level processors do not know types.

- * They treat all data as bits.
- * Operations apply to bits, but their result is meaningless if types do not match.

Type checking can be performed at different times.

- * *Dynamic typing*: check types at *run-time*; raise exception on a mismatch.
- * *Static typing*: check types before the program is run.
- * A type mismatch indicates an error which the programmer needs to correct.
- * Static typing allows the early detection of such errors.
- * Type information can also be used for optimisation.

A *type system* defines types and typing rules for a programming language.

- * Assembler: no types.
- * C: few types, uncontrolled type casts.
- * C++, Java: more types, type polymorphism.
- * Haskell: expressive type system, no type errors at run-time.
- * Python: dynamic typing.

3.3.2 Attribute Grammars

Simple type system for numbers:

The type system is implemented by an *attribute grammar*.

- * This is a CFG extended by attributes and rules.
- * An attribute stores information associated with non-terminals and terminals.
- * The rules describe how to calculate this information.

Attribute grammar for typing arithmetic expressions:

* attribute

* attribute

$$E = T$$
$$E = E A T$$
$$A = +$$
$$A = -$$
$$T = F$$
$$T = T M F$$
$$M = *$$
$$M = /$$
$$F = (E)$$
$$F = \text{num}$$
$$F = \text{id}$$

These are *synthesised* attributes.

- * The value on the left-hand side depends on the values of the right-hand side.
- * Information is propagated bottom-up in the syntax tree.
- * The type of a number constant is determined by the scanner.
- * The type of an identifier is determined by its declaration or by assignments.

Type checking uses the calculated attribute values.

- * The types of identifier and expression have to match in an assignment.
- * Arguments must be integers for integer division.
- * If a mismatch is found, a type error is reported or automatic conversions are added.

There are also *inherited* attributes.

- * The value for a non-terminal on the right-hand side depends on the other values.
- * Information is propagated top-down in the syntax tree.
- * In combination with synthesised attributes information is passed around in the syntax tree.
- * Attributes are often calculated during parsing without a separate pass over the tree.

3.4 Machine-Independent Optimisation

The fourth step of compiling is performed by the *machine-independent optimiser*.

- * Optimisation takes place on the syntax tree or another intermediate form.
- * It does not need to know the target processor.

The following example shows constant propagation, constant folding and dead code elimination.

- * Constants assigned to a variable can be propagated to uses of the variable and substituted.
- * Constant expressions can be evaluated and the result substituted (folded).
- * Code that does not affect the result can be removed.

```
x := 11;
y := 4;
z := 8 * x;
z := x * y;
z := z * 2 + x;
if z > 99 then
    z := 0
else
    z := 1
```

Optimisations often lead to further optimisations.

- * Constant folding enables constant propagation.
- * Constant propagation enables further constant folding.
- * Constant folding enables dead-code elimination.

Constant folding is the simplification of constant expressions at *compile-time*.

- * Evaluate arithmetic expressions that use only constants.
- * Values of variables are typically unknown, but might be known from preceding assignments.
- * Constants might also be inserted by a pre-processor.
- * Let ? denote an unknown value.



- * The evaluation must implement the semantics of the program that is compiled.
- * This may differ from operations of the machine on which the compiler runs.

Attribute grammar for evaluation of constant expressions:

* attribute

$$E = T$$
$$E = E A T$$
$$T = F$$
$$T = T M F$$
$$F = (E)$$
$$F = \text{num}$$
$$F = \text{id}$$

Compilers perform many kinds of optimisation leading to much faster running times:

- * move calculations out of loops;
- * reorder calculations;
- * inline code;
- * unroll loops;
- * eliminate common subexpressions;
- * remove tail-recursion.

3.5 Code Generation

The fifth step of compiling is performed by the *code generator*.

- * It traverses the abstract syntax tree or another intermediate form.
- * It emits code as soon as a sufficient portion is processed.
- * The output is a data structure representing code for a virtual machine.
- * It can be further analysed and optimised.
- * The program might also be translated to binary code for a real machine.
- * Thus compilation is *translation* from one language to another.

Programs have different parts, which are translated in different ways.

- * *Expressions* are program fragments that yield a value.
- * *Statements* are fragments that modify the values of variables.
- * *Declarations* provide information for type checking and other analyses.


3.5.1 A Virtual Machine

In the following, the target code runs on the Java virtual machine.

- * The machine features a *stack* to perform calculations.
- * Values of numbers and identifiers are pushed on the stack.
- * Operations remove their operands from the stack and push the result back.


Consider the expression $2 + 8 * 5$.

- * The code generated for this expression is:



- * `sipush`, `imul`, `iadd` are instructions of the Java virtual machine.
- * The instruction `sipush n` pushes the integer constant n on the stack.
- * The instruction `imul` pops two integers from the stack and pushes their product on the stack.
- * The instruction `iadd` similarly adds the two top integers.

* The steps of running this code are:



* Running the code generated for an expression leaves its value on top of the stack.

3.5.2 Code Generation for Expressions

Code can thus be generated recursively.

* Let $code(e)$ be the code generated for expression e .

* If n is a number, $code(n)$ is just

* $code(e_1 + e_2)$ is the sequence

* Similar code is generated for $e_1 - e_2$ and $e_1 * e_2$ and e_1 / e_2 , using `isub`, `imul` and `idiv`.

* The code generator is implemented by methods of the abstract syntax tree classes.

```
class Number_AST:
    def code(self):
        return 'sipush ' + self.number + '\n'

class Expression_AST:
    def code(self):
        op = { '+': 'iadd', '-': 'isub', '*': 'imul', '/': 'idiv' }
        return self.left.code() + self.right.code() + \
            op[self.op] + '\n'
```

* These methods return a string representation of the generated code.

* The resulting code is fed to a Java assembler to produce a Java class file.

* The Java class file can be executed using the Java virtual machine.

3.5.3 Identifiers in the Virtual Machine

Another kind of expression is an identifier.

- * Identifiers refer to variables.
- * Variables are like numbers: their value should be pushed on the stack.
- * Unlike constant numbers, the value of a variable can be changed by assignments.
- * Therefore the value of a variable is stored in the memory of the machine.
- * To change it, the *location* of that piece of storage is remembered.

Variables have two kinds of associated information:

A compiler needs to know the L-value.

- * The R-value can change, but the L-value is fixed at compile-time.
- * Obtain the R-value from the L-value by looking up the memory contents at that location.

A *frame* contains arguments, local data and the calculation stack of a method call.

- * Locations of *local variables* or statically allocated data are offsets relative to the frame.
- * Locations of *global variables* or dynamically allocated data are absolute memory addresses.
- * The following assumes just a single frame.

Consider the assignment $x := y + 2$.

- * Assume relative location 0 for *y* and 1 for *x*.
- * The code generated for this assignment is:

- * The instruction `iload n` pushes the value stored at relative location *n* on the stack.
- * The instruction `istore n` pops a value from the stack and stores it at relative location *n*.

3.5.4 Code Generation for Identifiers

During compilation, the *symbol table* keeps track of the location of each variable.

- * Upon encountering a new identifier, it is entered into the symbol table with a new location.
- * Relative locations are just consecutive numbers 0, 1, 2, ...

- * If the identifier is encountered subsequently, its location is looked up in the symbol table.

```
class Symbol_Table:
    def __init__(self):
        self.symbol_table = {}
    def location(self, identifier):
        '''Returns the location of an identifier.
        A new identifier is entered with a new location.'''
        if identifier in self.symbol_table:
            return self.symbol_table[identifier]
        index = len(self.symbol_table)
        self.symbol_table[identifier] = index
        return index
```

- * $code(v)$ for a variable v with location $l(v)$ is just

- * The following class implements code generation for identifiers:

```
class Identifier_AST:
    def code(self):
        loc = symbol_table.location(self.identifier)
        return 'iload ' + str(loc) + '\n'
```


Using relative locations 0 and 1 for y and x , the code generated for $y + 8 * x$ is:

3.5.5 Code Generation for Assignments

An assignment statement changes the value of a variable.

- * The left-hand side of an assignment is an identifier.
- * The right-hand side of an assignment is an expression.
- * The generated code evaluates the expression and stores its value at the variable's location.
- * $code(v := e)$ for expression e and variable v with location $l(v)$ is:

* The code generated for the assignment `x := y + 8 * x` is:



* The following class implements assignments and their code generation:

```
class Assign_AST:
    def __init__(self, identifier, expression):
        self.identifier = identifier
        self.expression = expression
    def code(self):
        loc = symbol_table.location(self.identifier.identifier)
        return self.expression.code() + 'istore ' + str(loc) + '\n'
```

* The function for parsing assignments is:

```
def assignment():
    value = consume(ID)[1]
    ident = Identifier_AST(value)
    consume(BEC)
    expr = expression()
    return Assign_AST(ident, expr)
```

* The token BEC represents `:=`.

3.5.6 Control flow in the Virtual Machine

Conditionals and loops involve a change of *control flow*.

* Virtual machine instructions are executed sequentially.

* This matches a sequence of assignments.

* For other statements, the executed code depends on conditions known only at run-time.

* This is implemented by *labels* and *jumps*.

* Labels mark positions in the virtual machine code.

```
11:
```

* An *unconditional jump* transfers control to a given label.

```
goto 11
```

* A *conditional jump* does the same, but only if a given condition holds.

```
if_icmpeq 11
```


- * `if_icmpeq l` pops the top two values from the stack, and jumps to `l` if they are equal.
- * If they are not equal, execution continues with the next instruction.

Consider the following program:

```
n := 2;
while n > 0 do
  n := n - 1
end
```

The code generated for this program is:

3.5.7 Code Generation for While-Loops

A condition compares two expressions.

- * The values of both expressions are calculated first.
- * The comparison is performed similarly to an arithmetic operation.
- * Unlike arithmetic, which leaves the result on the stack, a conditional jump is emitted.
- * $\text{false_code}(e_1 < e_2, l)$ is the sequence


- * The additional label `l` is the destination of the jump in case the comparison is *false*.

- * `if_icmpge` l jumps to l if $e_1 \geq e_2$, otherwise continues with the next instruction.
- * `if_icmpne`, `if_icmple`, `if_icmpgt`, `if_icmpeq`, `if_icmplt` are similar comparisons.
- * The following class implements comparisons and their code generation:

```
class Comparison_AST:
    def __init__(self, left, op, right):
        self.left = left
        self.op = op
        self.right = right
    def false_code(self, label):
        op = { '<': 'if_icmpge', '<=': 'if_icmpgt',
              '=': 'if_icmpne', '!=': 'if_icmpeq',
              '>': 'if_icmple', '>=': 'if_icmplt' }
        return self.left.code() + self.right.code() + \
            op[self.op] + ' ' + label + '\n'
```

A while-loop has a condition and a body, which is a sequence of statements.

- * The condition is evaluated first.
- * If it is true, the body is executed; this needs a conditional jump.
- * After execution of the body, the above is repeated; this needs an unconditional jump.
- * `code(while c do s end)` is the sequence



- * This uses two new labels l_1 and l_2 .
- * The following class generates code for while-loops:

```
class While_AST:
    def __init__(self, condition, body):
        self.condition = condition
        self.body = body
    def code(self):
        l1 = new_label()
        l2 = new_label()
        return l1 + ':\n' + self.condition.false_code(l2) + \
            self.body.code() + 'goto ' + l1 + '\n' + l2 + ':\n'
```

- * `new_label()` returns the string representation of a new label.

3.6 Machine-Dependent Optimisation

The sixth step of compiling is performed by the *machine-dependent optimiser*.

- * It modifies the generated code.
- * The optimisations are specific for each target processor.
- * *Peephole optimisation* looks at a small part of generated instructions.
- * They might be replaced with a shorter or faster sequence of instructions.

Consider the following fragment of a Java program:

```
int x, y, z;  
z = 7;  
x = y = z = z + 1;
```

The unoptimised code is:

```
sipush 7  
istore 2  
iload 2  
sipush 1  
iadd  
dup      -- duplicate top element of stack  
istore 2  
dup  
istore 1  
dup  
istore 0  
pop      -- remove top element of stack
```

Realistic machines are more detailed than discussed above.

- * They have stores for program, stack, heap.
- * Their stacks are divided into frames, each of which belongs to a method call.
- * They have registers such as program counter, stack pointer, frame pointer.
- * The underlying interpreter loop fetches, decodes and executes instructions.
- * The benefit of a virtual machine is that it abstracts from these specific details.
- * The virtual machine can be implemented on a variety of platforms.
- * This is an instance of adding a *level of indirection*.