

COSC264

Introduction to Computer Networks and the Internet

Transport Layer Protocols: TCP Reliable Data Transfer

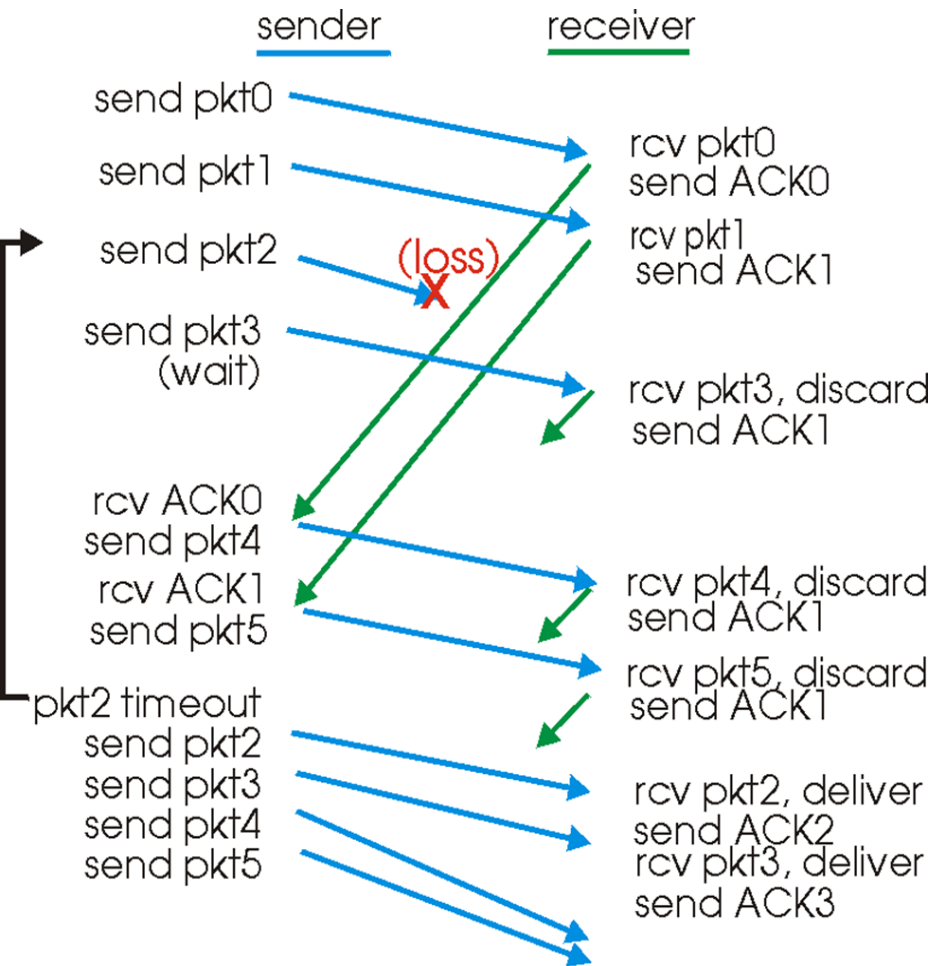
Dr. Barry Wu

Wireless Research Centre

University of Canterbury

barry.wu@canterbury.ac.nz

Go-Back-N: an example



0 1 2 3 4 5 6 7

0 1 2 3 4 5 6 7

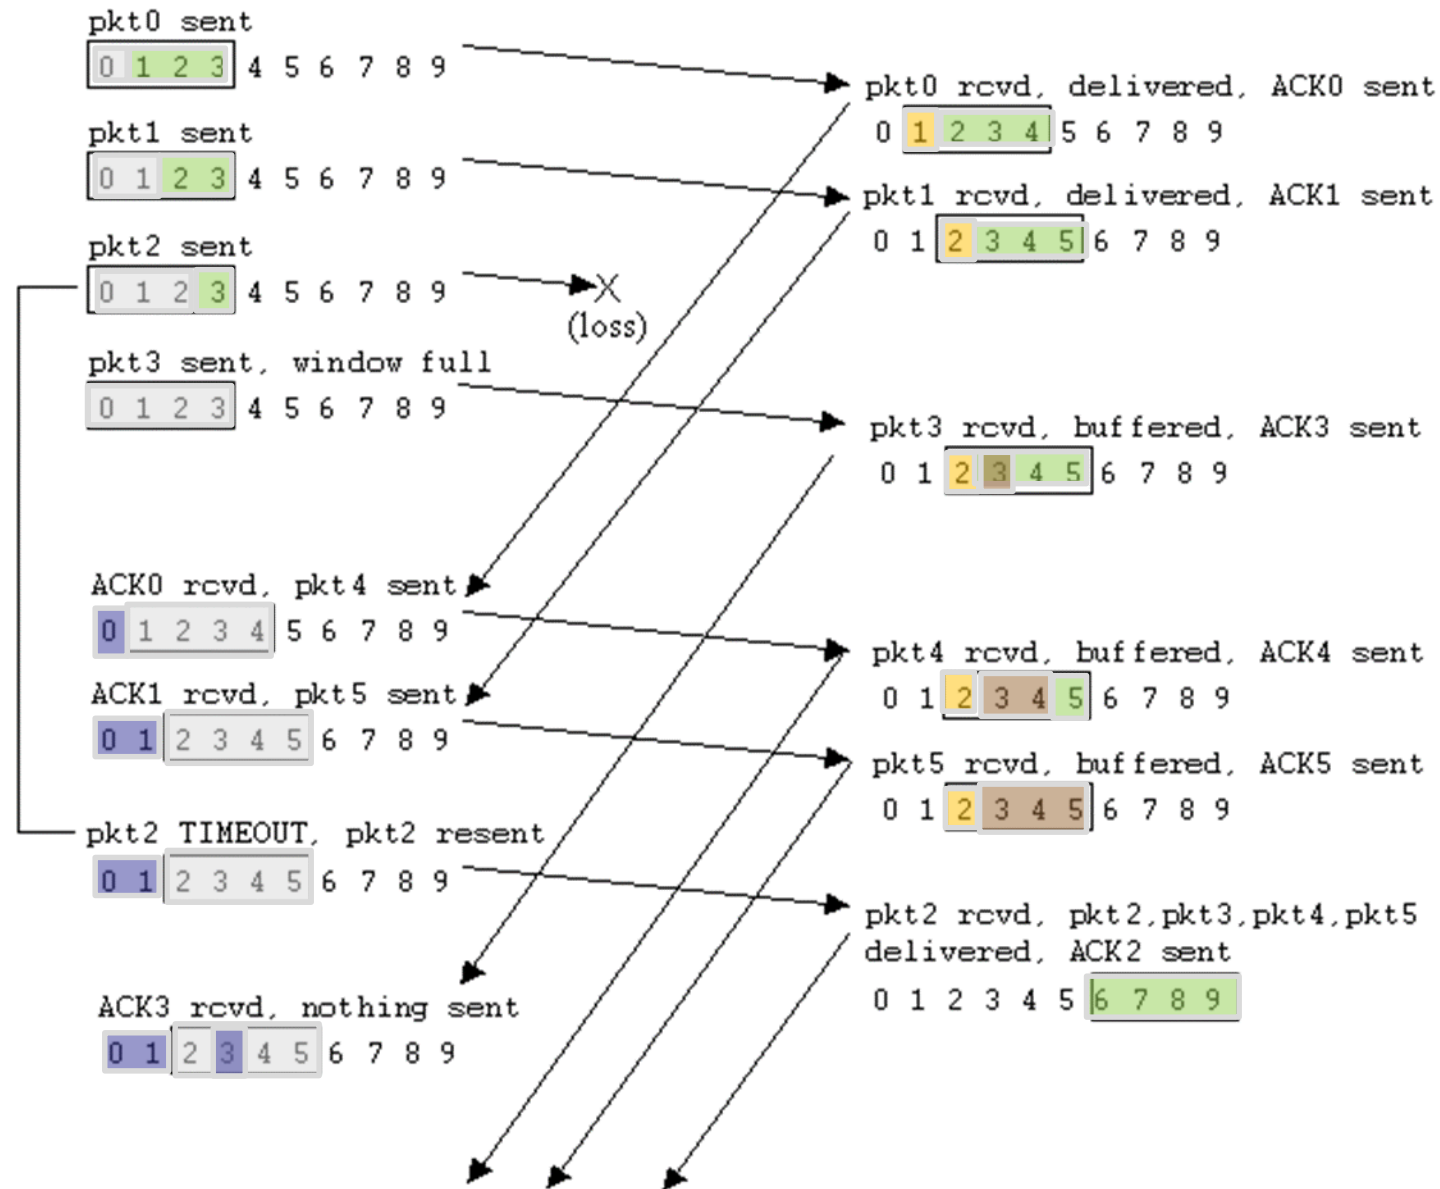
0 1 2 3 4 5 6 7 0

0 1 2 3 4 5 6 7 0 1

0 1 2 3 4 5 6 7 0 1

[AK05, KR3]


SR: an example (after-receipt window state)



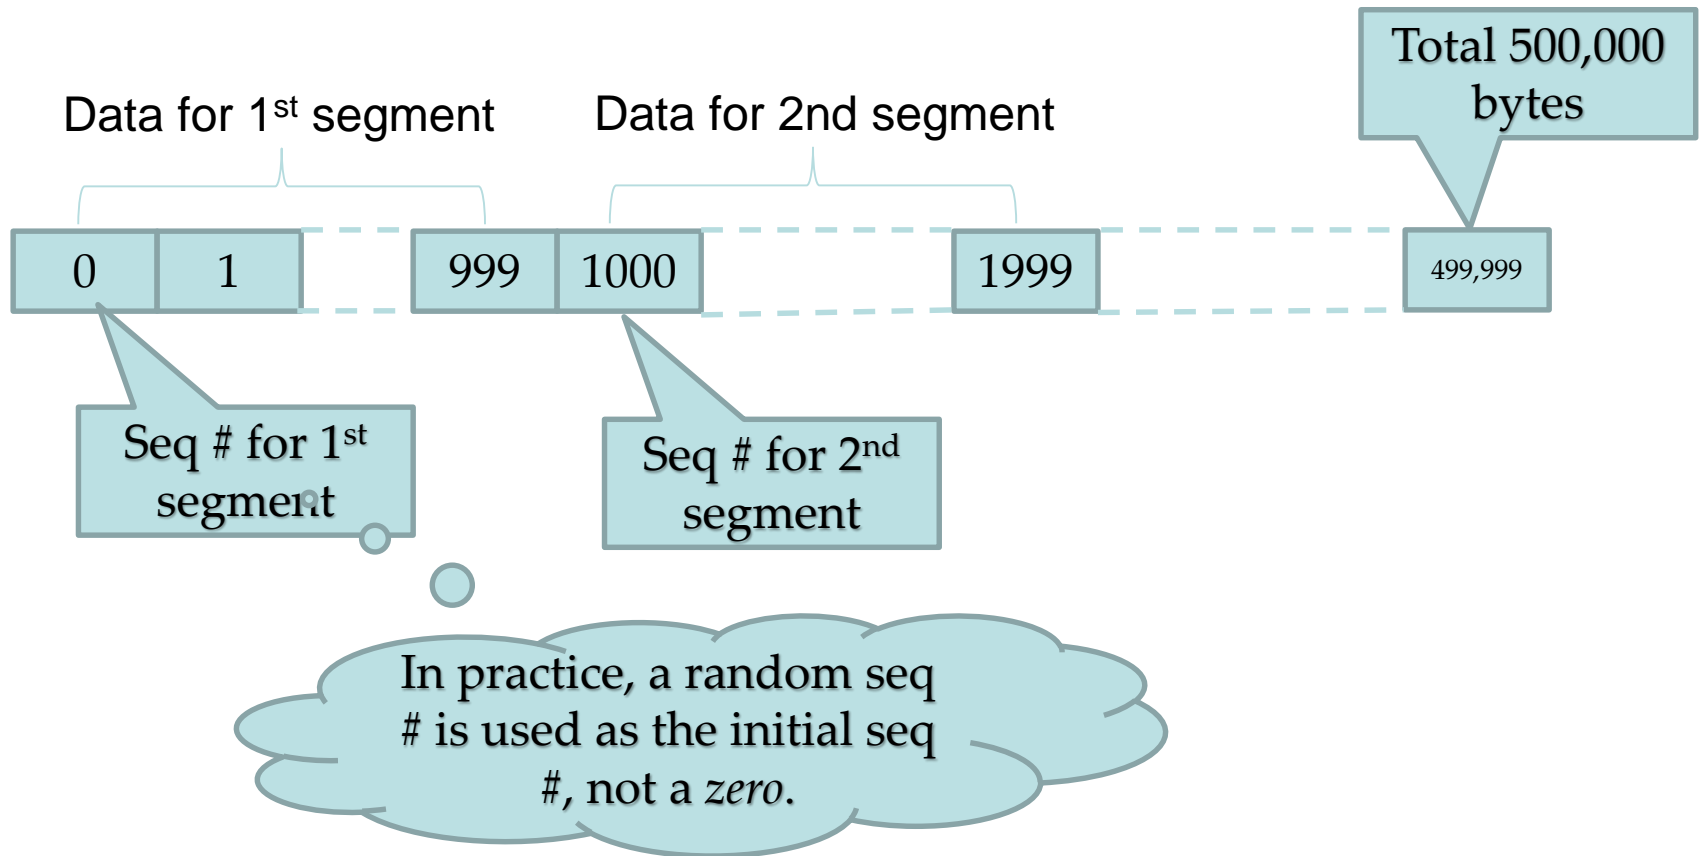
Outline


- TCP Reliable data transfer

TCP basics





TCP segments







Window of size N

 ACK'd seq#

 Usable seq#, not yet sent

 Sent seq#, not yet ACK'd

 Not usable seq#

←
Older packets

NextSeqNum = InitialSeqNum
SendBase = InitialSeqNum


TCP

sender (simplified)

```
loop (forever) {  
  switch(event)  
  event: data received from application above  
    create TCP segment with sequence number NextSeqNum  
    if (timer currently not running)  
      start timer  
    pass segment to IP  
    NextSeqNum = NextSeqNum + length(data)
```

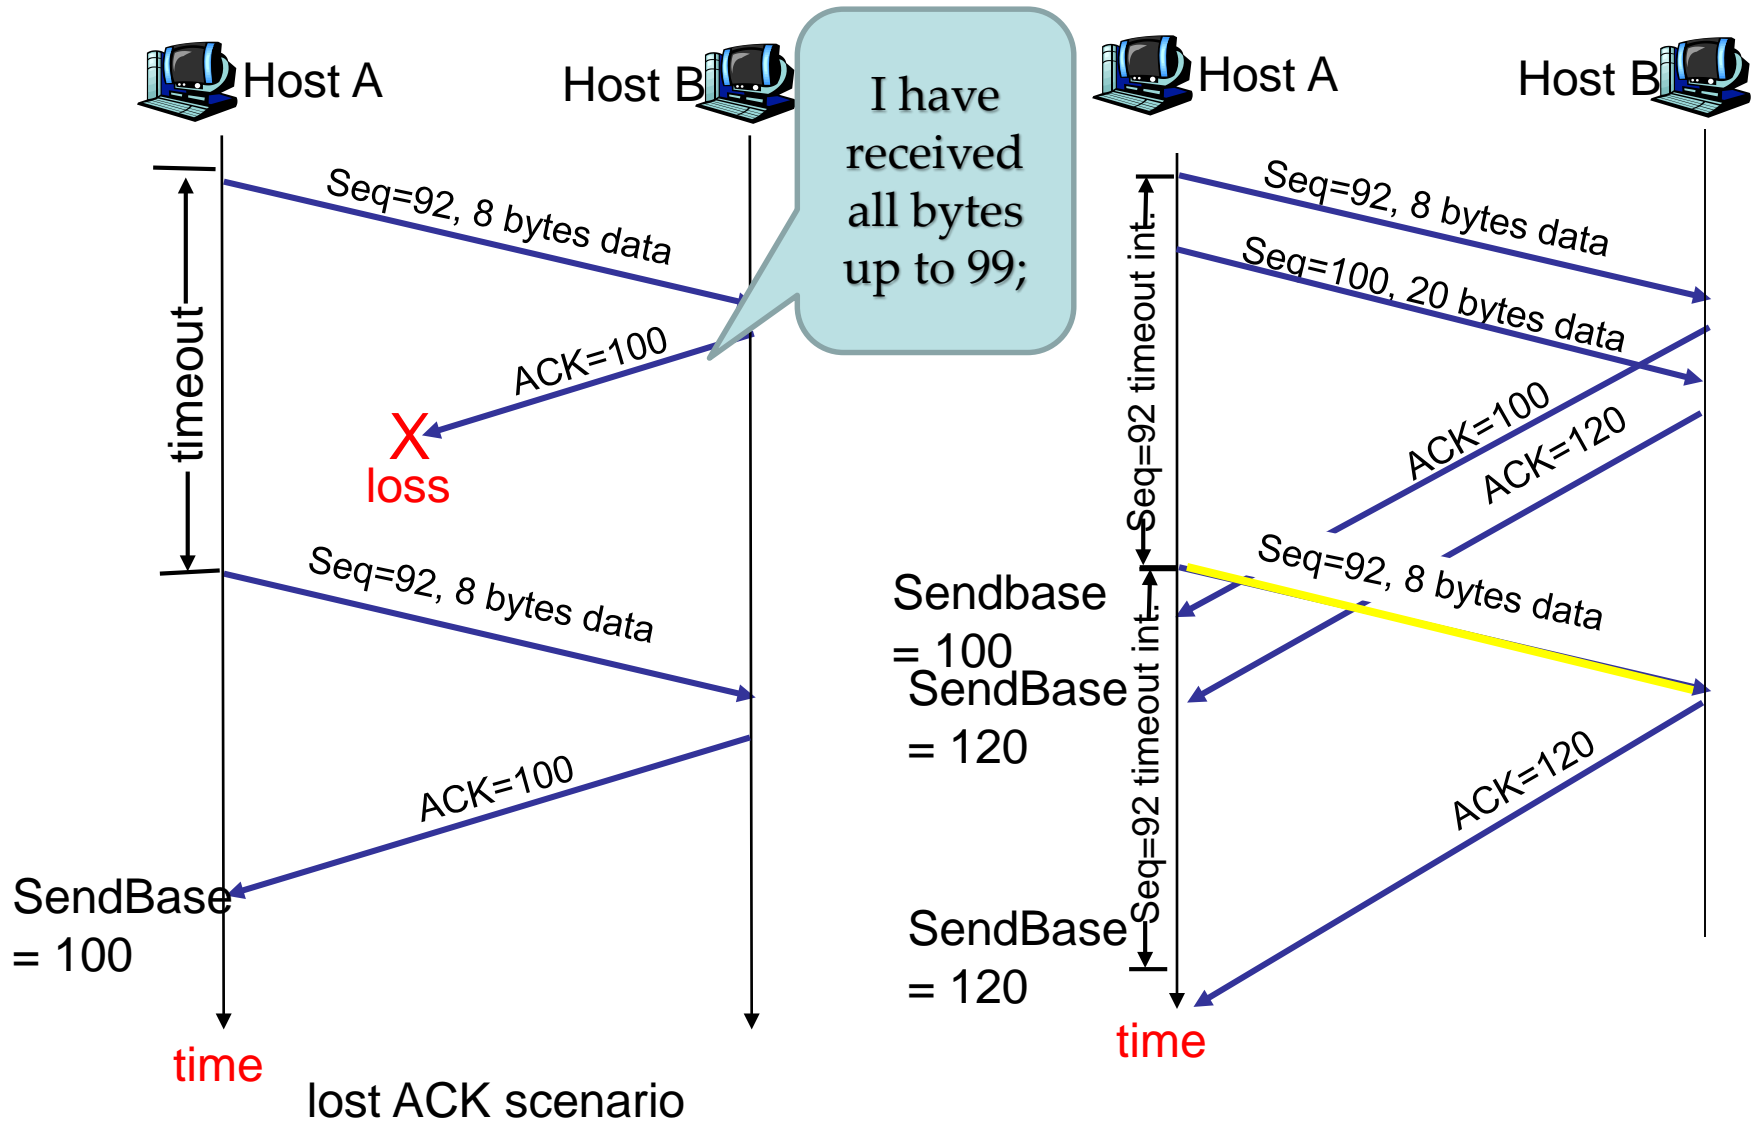
```
  event: timer timeout  
    retransmit not-yet-acknowledged segment with  
      smallest sequence number  
    start timer
```

```
  event: ACK received, with ACK field value of y  
    if (y > SendBase) {  
      SendBase = y  
      if (there are currently not-yet-acknowledged segments)  
        start timer  
    }  
}
```




```
} /* end of loop forever */
```


TCP: retransmission scenarios



TCP retransmission scenarios (more)



Cumulative ACK scenario

TCP sender events:

data rcvd from app:

- Create segment with seq #
- seq # is byte-stream number of first data byte in segment
- start timer if not already running (think of timer as for oldest unacked segment)

TCP sender events:

timeout:

- retransmit segment with the *smallest* seq#;
- restart timer


Ack rcvd:

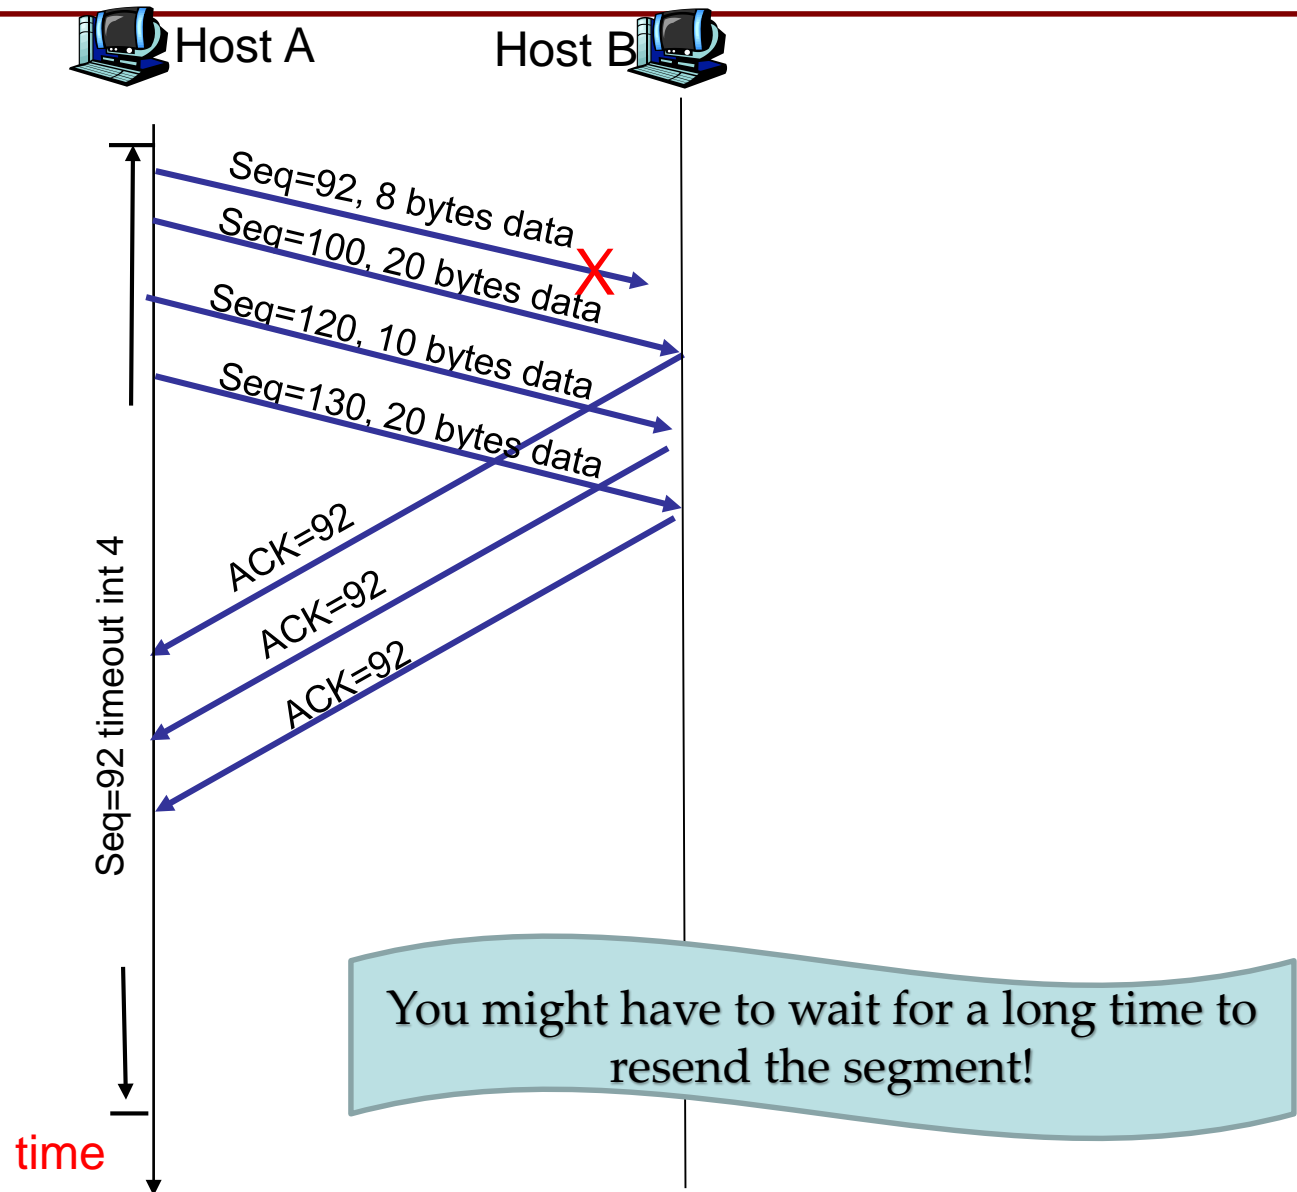
- Cumulative ACK;
- Slide window;
- ... (later);

Some refinement

timeout:

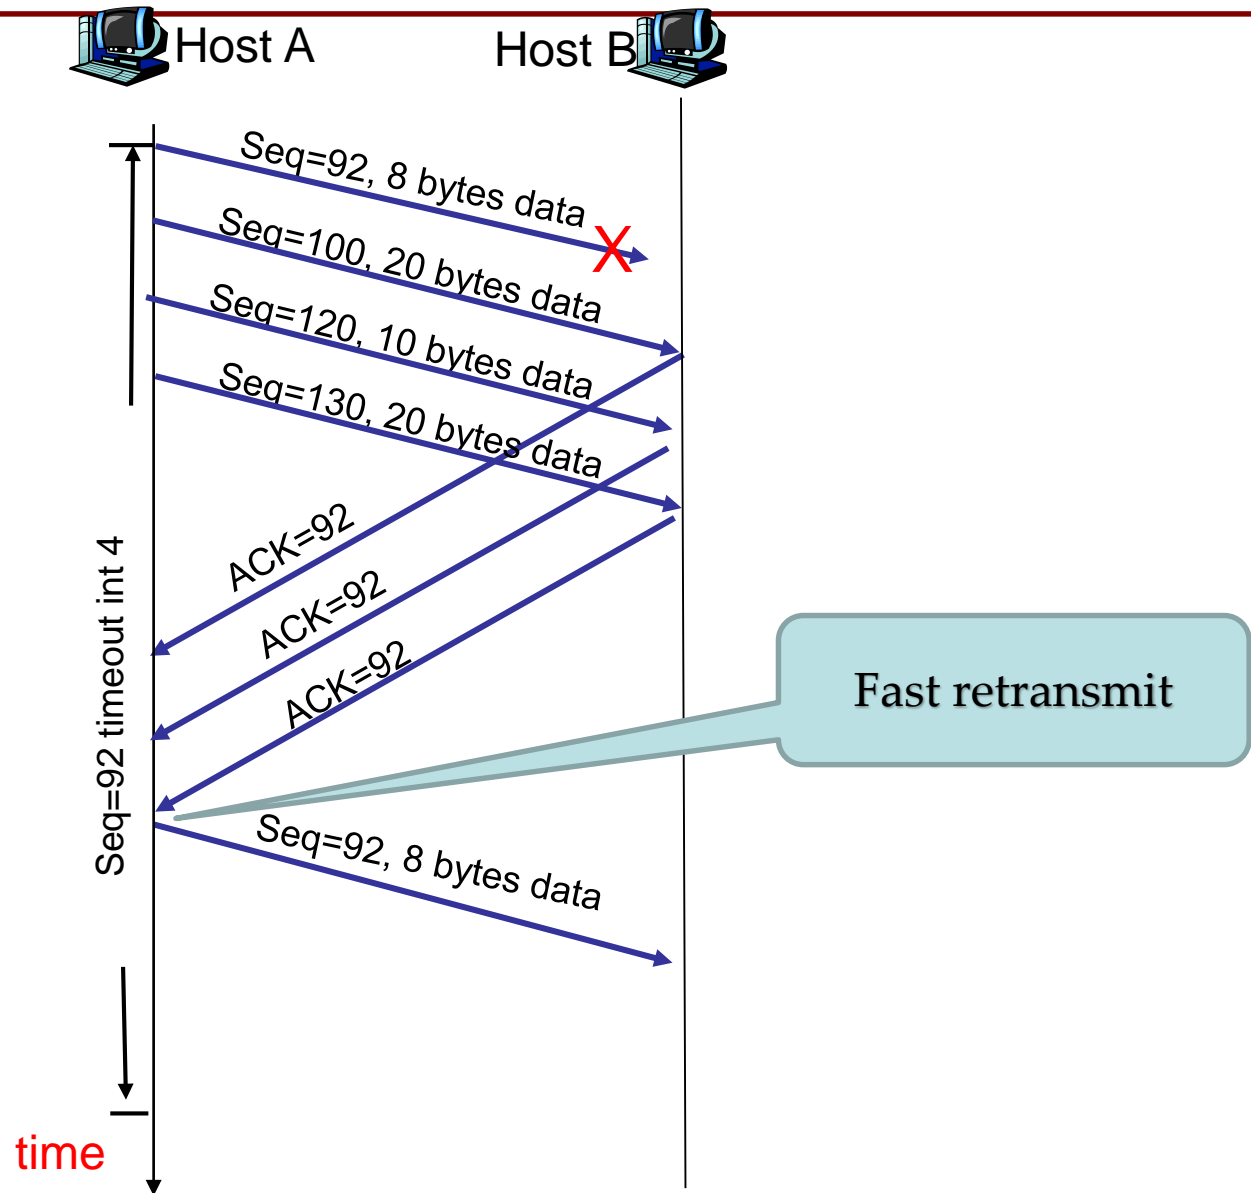
- ❑ When timeout expires after `TimeoutInterval` ;
 - Retransmits the unAck'd segment with smallest seq#;
(*not all of the unAck'd pkts, different from Go-Back-N*)
 - (new) `TimeoutInterval` = $2 \times$ (previous) `TimeoutInterval`
- ❑ Motivation: no packet was received for the last timeout period
 - Could be a congestion... *trying to slow down the sending process*;
- ❑ `TimeoutInterval` uses *initial value* when timer is started in events “call from above” or “receipt of ACK” is activated;





Fast Retransmit

- If sender receives 3 ACKs for the same data, it supposes that segment after the ACKed data was lost:
 - fast retransmit: resend segment before timer expires
- Time-out period is often relatively long:
 - long delay before resending lost packet
- Detect lost segments via duplicate ACKs.
 - Sender often sends many segments back-to-back
 - If segment is lost, there will likely be many duplicate ACKs.



Fast retransmit algorithm:

```
event: ACK received, with ACK field value of y
    if (y > SendBase) {
        SendBase = y
        if (there are currently not-yet-acknowledged segments)
            start timer
    }
    else {
        increment count of dup ACKs received for y
        if (count of dup ACKs received for y = 3) {
            resend segment with sequence number y
        }
    }
```

a duplicate ACK for
already ACKed segment

fast retransmit

After receiving 3 dup ACKs, the sender is very sure
that the related pkt has been lost;

TCP ACK generation [RFC 1122, RFC 2581]

Event at Receiver	TCP Receiver action
Arrival of in-order segment with expected seq #. All data up to expected seq # already ACKed	Delayed ACK. Wait up to 500ms for next segment. If no next segment, send ACK
Arrival of in-order segment with expected seq #. One other segment has ACK pending	Immediately send single cumulative ACK, ACKing both in-order segments
Arrival of out-of-order segment higher-than-expect seq. # . Gap detected	Immediately send duplicate ACK, indicating seq. # of next expected byte
Arrival of segment that partially or completely fills gap	Immediately send ACK, provided that segment starts at lower end of gap

GBN or SR?

TCP rdt

- Cumulative ACK (similar to GBN)
- Many TCP implementation buffers out-of-order segments (similar to SR);
- When there is timeout; sender retransmits the unAck'ed segments with smallest seq# only! (no “go-back-n”)
- A mixture;

TCP reliable data transfer: a summary

- TCP creates rdt service on top of IP's unreliable service
- Pipelined segments
- Cumulative ACKs
- TCP uses single retransmission timer
- Retransmissions are triggered by:
 - timeout events
 - duplicate ACKs (*fast retransmit*)

Tricks for reliable data transfer

The length of timeout intervals?

Mechanism	Use, Comments
Checksum	To detect bit errors
Timer	To timeout/retransmit a packet (lost/premature packet)
Sequence number	To detect a lost packet (gap in seq#) To detect a duplicate packet (duplicate seq#)
Acknowledgement	To notify successful reception (individual/cumulative)
Negative ACK	To notify unsuccessful reception
Window, pipelining	To improve sender utilisation (utilisation vs performance)

TCP Round Trip Time and Timeout

Q: how to set TCP timeout value?

- longer than RTT
 - but RTT varies
- too short: premature timeout
 - unnecessary retransmissions
- too long: slow reaction to segment loss

TCP Round Trip Time and Timeout

Q: how to estimate RTT?

- **SampleRTT**: measured time from segment transmission until ACK receipt
 - ignore retransmissions (*do not sample retransmitted segments*);
- **SampleRTT** will vary, want estimated RTT “typical”
 - average several recent measurements

TCP Round Trip Time and Timeout

$$\text{EstimatedRTT} = (1 - \alpha) * \text{EstimatedRTT} + \alpha * \text{SampleRTT}$$

$$E_0 = S_0$$

$$E_1 = (1 - \alpha)E_0 + \alpha S_1 = (1 - \alpha)S_0 + \alpha S_1$$


$$E_2 = (1 - \alpha)E_1 + \alpha S_2 = (1 - \alpha)^2 S_0 + \alpha(1 - \alpha)S_1 + \alpha S_2$$

$$E_n = (1 - \alpha)E_{n-1} + \alpha S_n = (1 - \alpha)^n S_0 + \alpha(1 - \alpha)^{n-1} S_1 + \alpha(1 - \alpha)^{n-2} S_2 + \dots + \alpha(1 - \alpha) S_{n-1} + \alpha S_n$$

- ❑ *Exponential weighted moving average (EWMA)*
- ❑ influence of past sample decreases exponentially fast
- ❑ typical value: $\alpha = 0.125$
- ❑ The weighted average puts more weight on recent samples which better reflect the current congestion

Example RTT estimation:

RTT: gaia.cs.umass.edu to fantasia.eurecom.fr



TCP Round Trip Time and Timeout

Setting the timeout

- To measure the variability of the RTT;
- first estimate of how much SampleRTT deviates from EstimatedRTT:

$$\text{DevRTT} = (1-\beta) * \text{DevRTT} + \beta * |\text{SampleRTT} - \text{EstimatedRTT}|$$

(typically, $\beta = 0.25$)

Then set timeout interval:

$$\text{TimeoutInterval} = \text{EstimatedRTT} + 4 * \text{DevRTT}$$


Tricks for reliable data transfer

The length of timeout intervals?

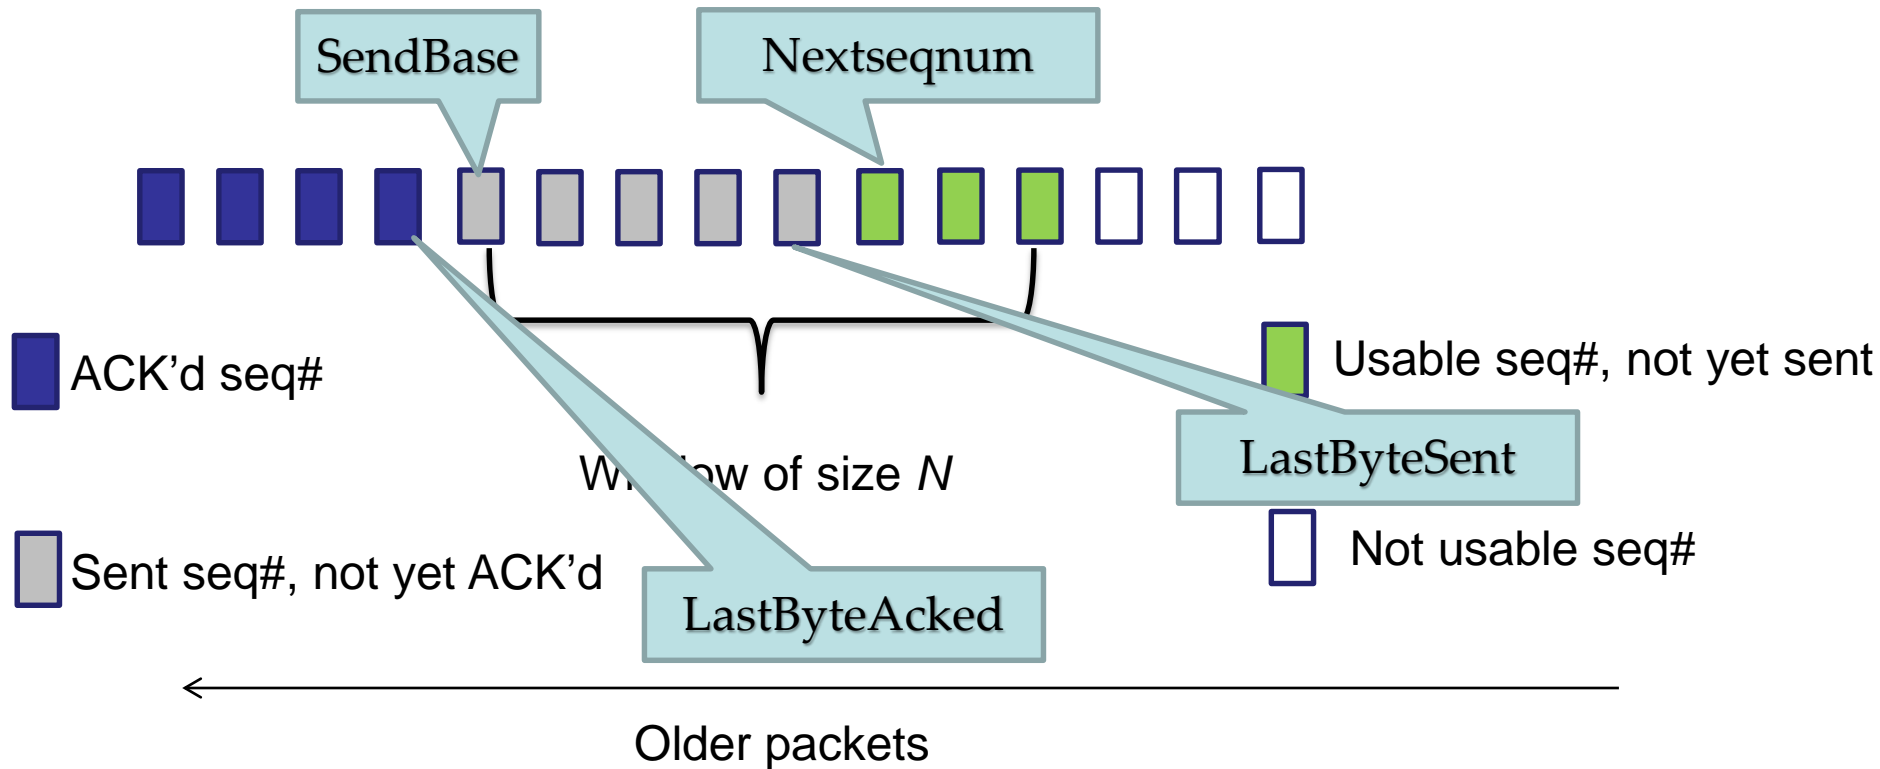
Mechanism	Use, Comments
Checksum	To detect bit errors
Timer	To timeout/retransmit a packet (lost/premature packet)
Sequence number	To detect a lost packet (gap in seq#) To detect a duplicate packet (duplicate seq#)
Acknowledgement	To notify successful reception (individual/cumulative)
Negative ACK	To notify unsuccessful reception
Window, pipelining	To improve sender utilisation (utilisation vs performance)

How to set the window size?

The journey of a packet



Problems	Causes	Solutions
Bit error	e.g., signal attenuation/noise	Error detection and correction
Buffer overflow	e.g., Speed-mismatch; Too much traffic;	Flow control and congestion control
Lost packet	e.g., buffer overflow at host/router	Acknowledgement and retransmission (ARQ) - RDT
Out of order	e.g. an early packet gets lost and retransmitted; a later one arrives first.	Acknowledgement and retransmission (ARQ) - RDT



$$\text{LastByteSent} - \text{LastByteAked} \leq \min\{\text{CongWin}, \text{RcvWindow}\};$$
 CongWin for congestion control; RcvWindow for flow control;

References

- [KR3] James F. Kurose, Keith W. Ross, *Computer networking: a top-down approach featuring the Internet*, 3rd edition.
- [PD5] Larry L. Peterson, Bruce S. Davie, *Computer networks: a systems approach*, 5th edition
- [TW5] Andrew S. Tanenbaum, David J. Wetherall, *Computer network*, 5th edition
- [LHBi]Y-D. Lin, R-H. Hwang, F. Baker, *Computer network: an open source approach*, International edition

Acknowledgements

- All slides are developed based on slides from the following two sources:
 - Dr DongSeong Kim's slides for COSC264, University of Canterbury;
 - Prof Aleksandar Kuzmanovic's lecture notes for CS340, Northwestern University,
https://users.cs.northwestern.edu/~akuzma/classes/CS340-w05/lecture_notes.htm