

COSC264

Introduction to Computer Networks and the Internet

Transport Layer Protocols: UDP and TCP

Dr. Barry Wu

Wireless Research Centre

University of Canterbury

barry.wu@canterbury.ac.nz

Outline

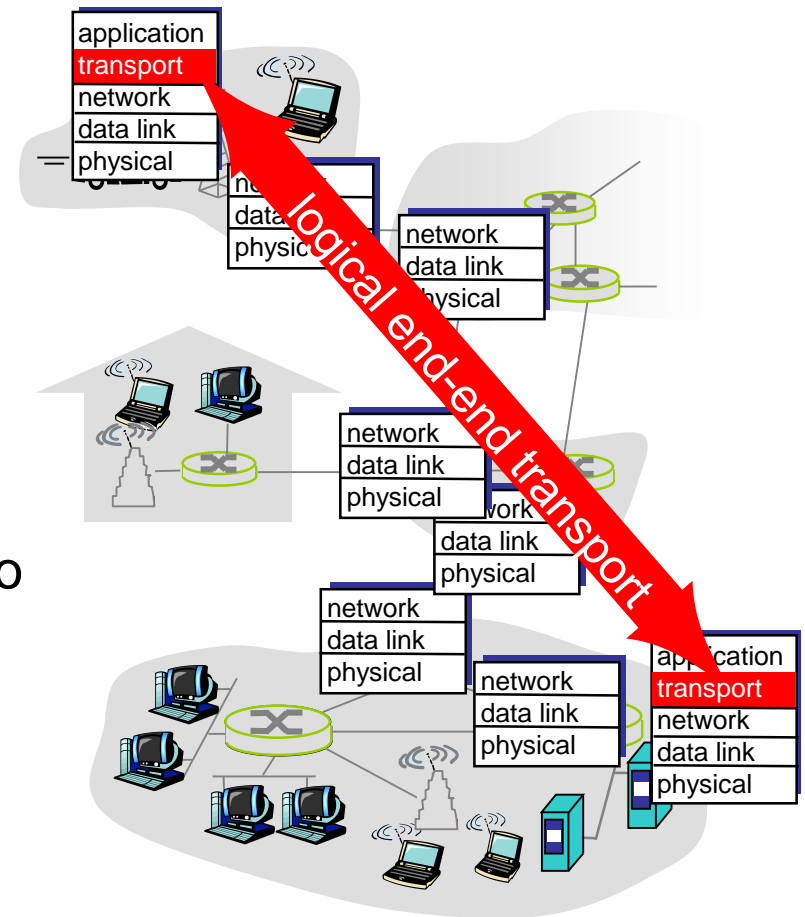
- Transport-layer services
- Multiplexing and demultiplexing
- Connectionless transport: UDP
- Connection-oriented transport: TCP

Popular Internet Applications and their underlying transport protocols

Application	Application-layer protocol	Underlying transport protocol
Email	SMTP	TCP
Web	HTTP	TCP
Routing protocol	BGP	TCP
Routing protocol	RIP	Typically UDP
Name translation	DNS	Typically UDP
Streaming multimedia	Typically proprietary	Typically UDP
Internet telephony	Typically proprietary	Typically UDP

Transport Protocols

- Provide *logical communication* between application processes running on different hosts
- Run on end hosts
 - Sender side: breaks application messages into **segments**, and passes to network layer
 - Receiver side: reassembles segments into messages, passes to application layer
- Multiple transport protocols available to applications
 - Internet: TCP and UDP
 - Other: RSVP (The Resource Reservation Protocol (RSVP))



Services provided by the Internet transport layer

- Extending host-to-host delivery to process-to-process delivery;
 - Transport-layer multiplexing and demultiplexing;
- Reliable data transfer;
- Flow control;
- Congestion control;

A host has many processes, communicating to many different processes at different remote hosts;
The Internet is a shared channel (with packet switching)!

Different destination hosts – network layer

Different processes -- transport layer

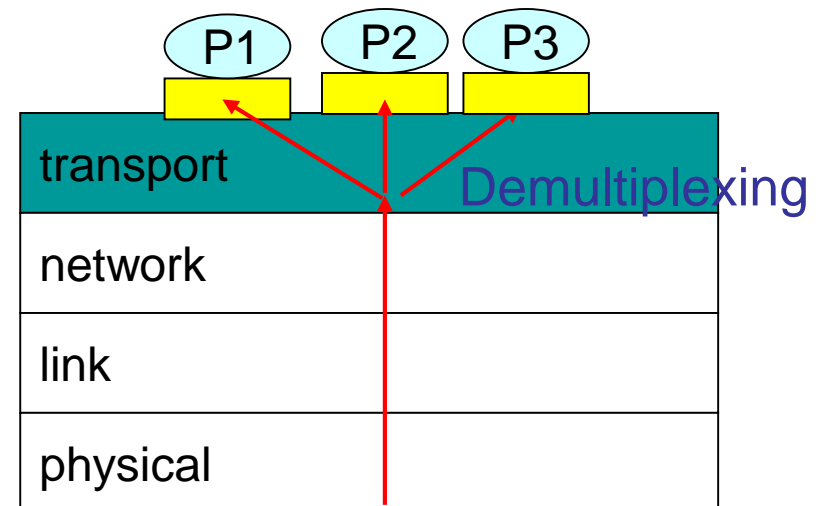
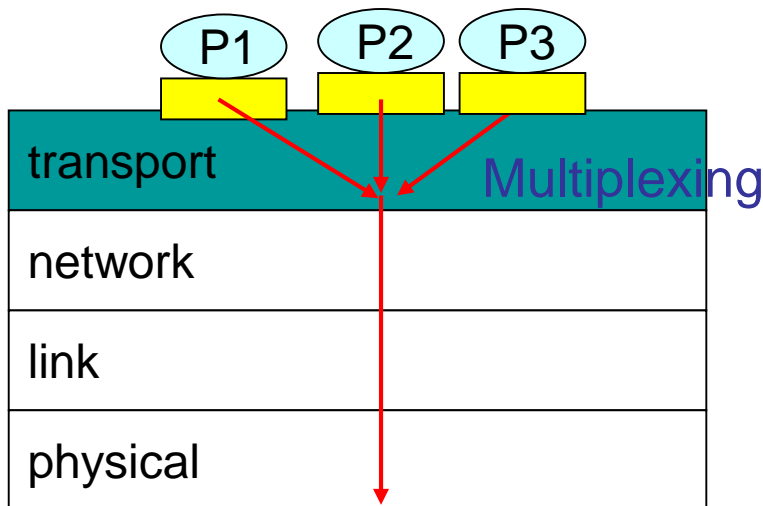
A shared channel – multiplexing and demultiplexing

Outline


- Transport-layer services
- Multiplexing and demultiplexing
- Connectionless transport: UDP
- Connection-oriented transport: TCP

Multiplexing/demultiplexing – Definition


- Demultiplexing
 - To deliver the data in a transport-layer segment to the correct socket
- Multiplexing
 - To create segments and pass them to the network layer.



How demultiplexing works



- **host receives IP datagrams**
 - each datagram has source IP address, destination IP address
 - each datagram carries 1 transport-layer segment
- each segment has source, destination port number
- **host uses IP addresses & port numbers to direct segment to appropriate socket**

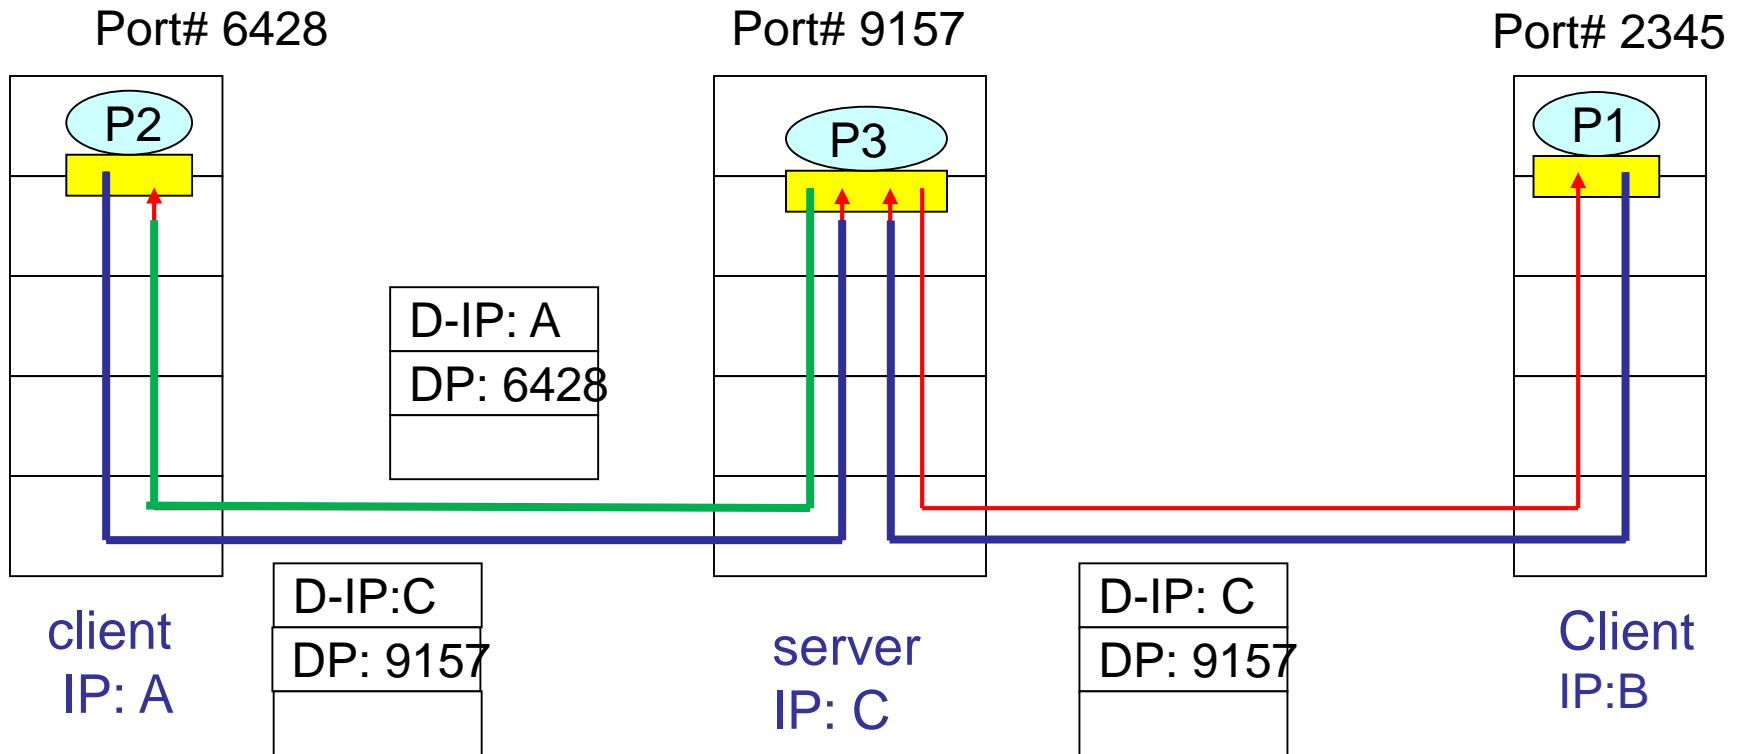


TCP/UDP segment format

Connectionless demultiplexing

- UDP socket identified by two-tuple:
(dest IP address, dest port number)
- When host receives UDP segment:
 - checks destination port number in segment
 - directs UDP segment to socket with that port number
- Packets with different source IP addresses and/or source port numbers, but with the same destination IP and port number, will be directed to same socket.

Connectionless demux (cont)

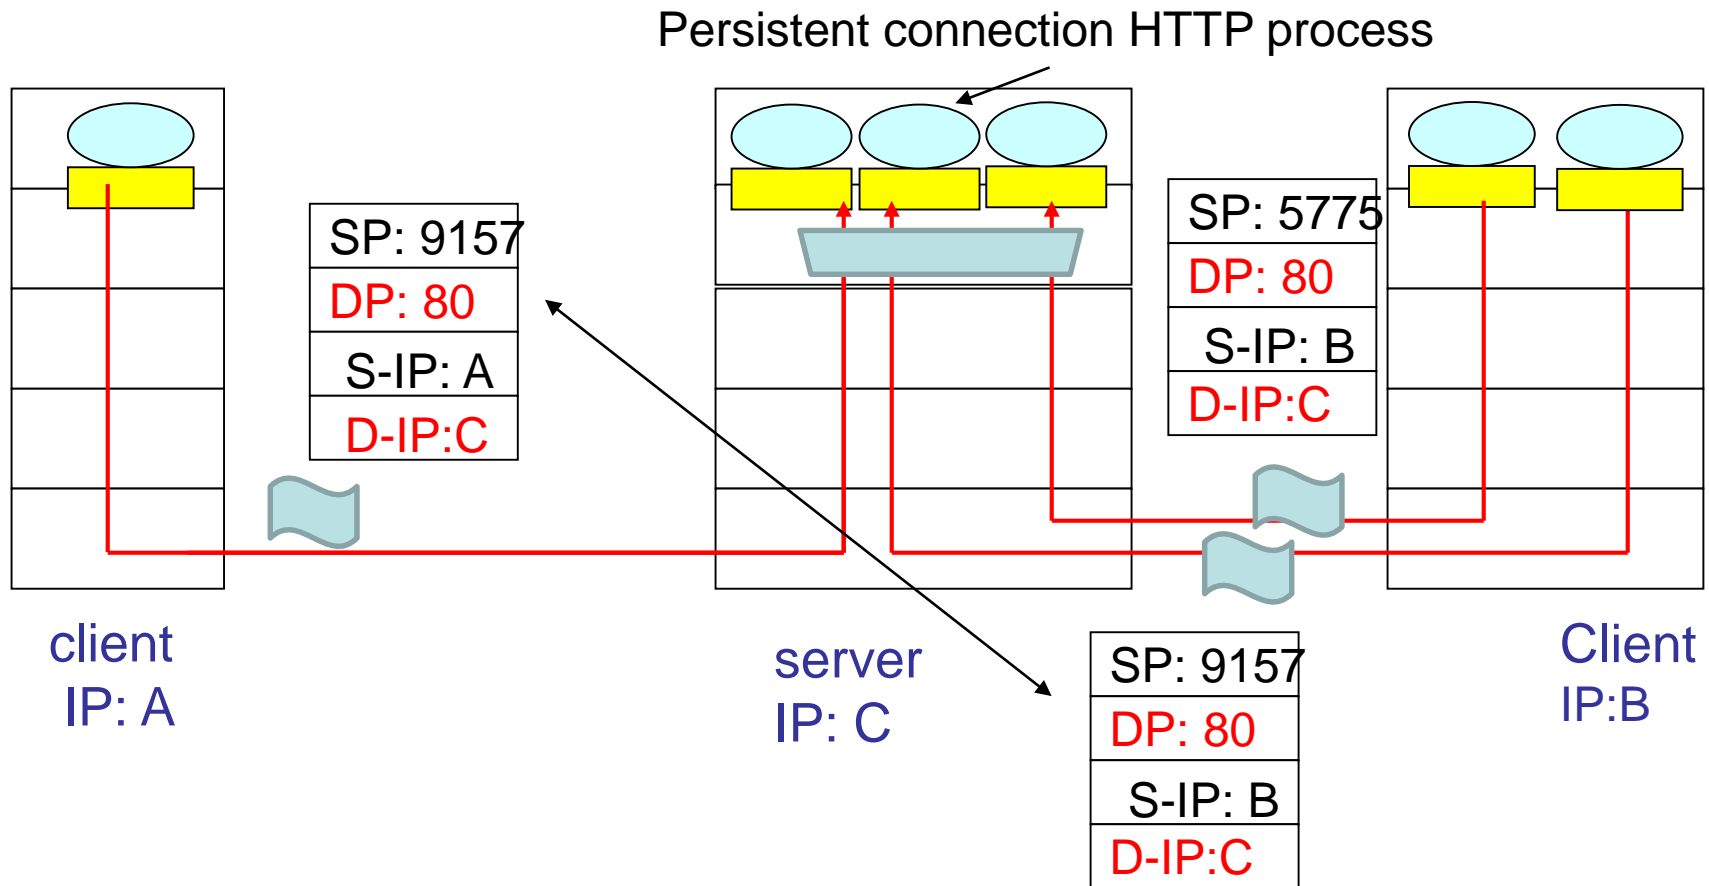


Source port provides "return address"

Connection-oriented demux

- TCP socket identified by 4-tuple:
 - source IP address
 - source port number
 - dest IP address
 - dest port number
- recv host uses *all four values* to direct segment to appropriate socket
- Server host may support many simultaneous TCP sockets:
 - each socket identified by its own 4-tuple
- Web servers have different sockets for each connecting client

Connection-oriented demux (cont)



The same Web server has different sockets for different processes at clients.

Outline

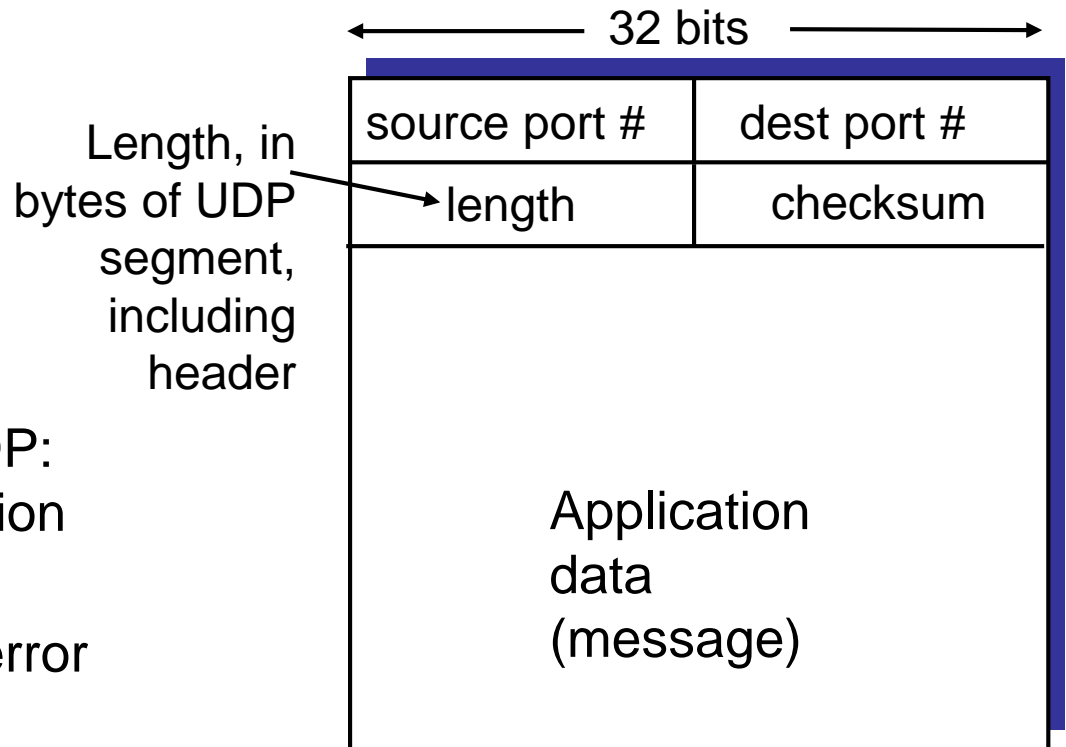
- Transport-layer services
- Multiplexing and demultiplexing
- Connectionless transport: UDP
- Connection-oriented transport: TCP

UDP: User Datagram Protocol [RFC 768]

- “no frills,” “bare bones” Internet transport protocol
- “best effort” service, UDP segments may be:
 - lost
 - delivered out of order to app
- *connectionless*:
 - no handshaking between UDP sender, receiver
 - each UDP segment handled independently of others

UDP: more

- often used for streaming multimedia apps
 - loss tolerant
 - rate sensitive
- other UDP uses
 - RIP, DNS, SNMP
- reliable transfer over UDP: add reliability at application layer
 - application-specific error recovery!



UDP segment format

A UDP segment captured by Wireshark

12	0.508950	10.34.40.169	132.181.2.225	DNS
13	0.509041	132.181.2.225	10.34.40.169	DNS
14	0.509462	10.34.40.169	132.181.2.225	DNS
15	0.510280	132.181.2.225	10.34.40.169	DNS
16	0.510354	132.181.2.225	10.34.40.169	DNS
17	33.503303	10.34.40.169	132.181.2.225	DNS
18	33.504666	132.181.2.225	10.34.40.169	DNS
19	61.002465	10.34.40.169	132.181.2.225	DNS
20	61.004045	132.181.2.225	10.34.40.169	DNS
21	61.013117	10.34.40.169	132.181.2.225	DNS
22	61.014578	132.181.2.225	10.34.40.169	DNS

- ▷ Frame 12: 88 bytes on wire (704 bits), 88 bytes captured (704 bits) on interface 0
- ▷ Ethernet II, Src: IntelCor_b6:fe:63 (80:19:34:b6:fe:63), Dst: JuniperN_ef:61:00 (2c:21:31:ef:61:00)
- ▷ Internet Protocol Version 4, Src: 10.34.40.169, Dst: 132.181.2.225
- ◀ User Datagram Protocol, Src Port: 63507, Dst Port: 53
 - Source Port: 63507
 - Destination Port: 53
 - Length: 54
 - Checksum: 0xe576 [unverified]
 - [Checksum Status: Unverified]
 - [Stream index: 6]
- ▷ Domain Name System (query)

Why is there a UDP?

- no connection establishment (which can add delay)
- simple: no connection state at sender, receiver
- small segment header (8 bytes)
- no congestion control: UDP can blast away as fast as desired

Outline

- Transport-layer services
- Multiplexing and demultiplexing
- Connectionless transport: UDP
- Connection-oriented transport: TCP


Outline

- Transport-layer services
- Multiplexing and demultiplexing
- Connectionless transport: UDP
- **Connection-oriented transport: TCP**
 - Overview and segment structure
 - Connection management

TCP: Overview

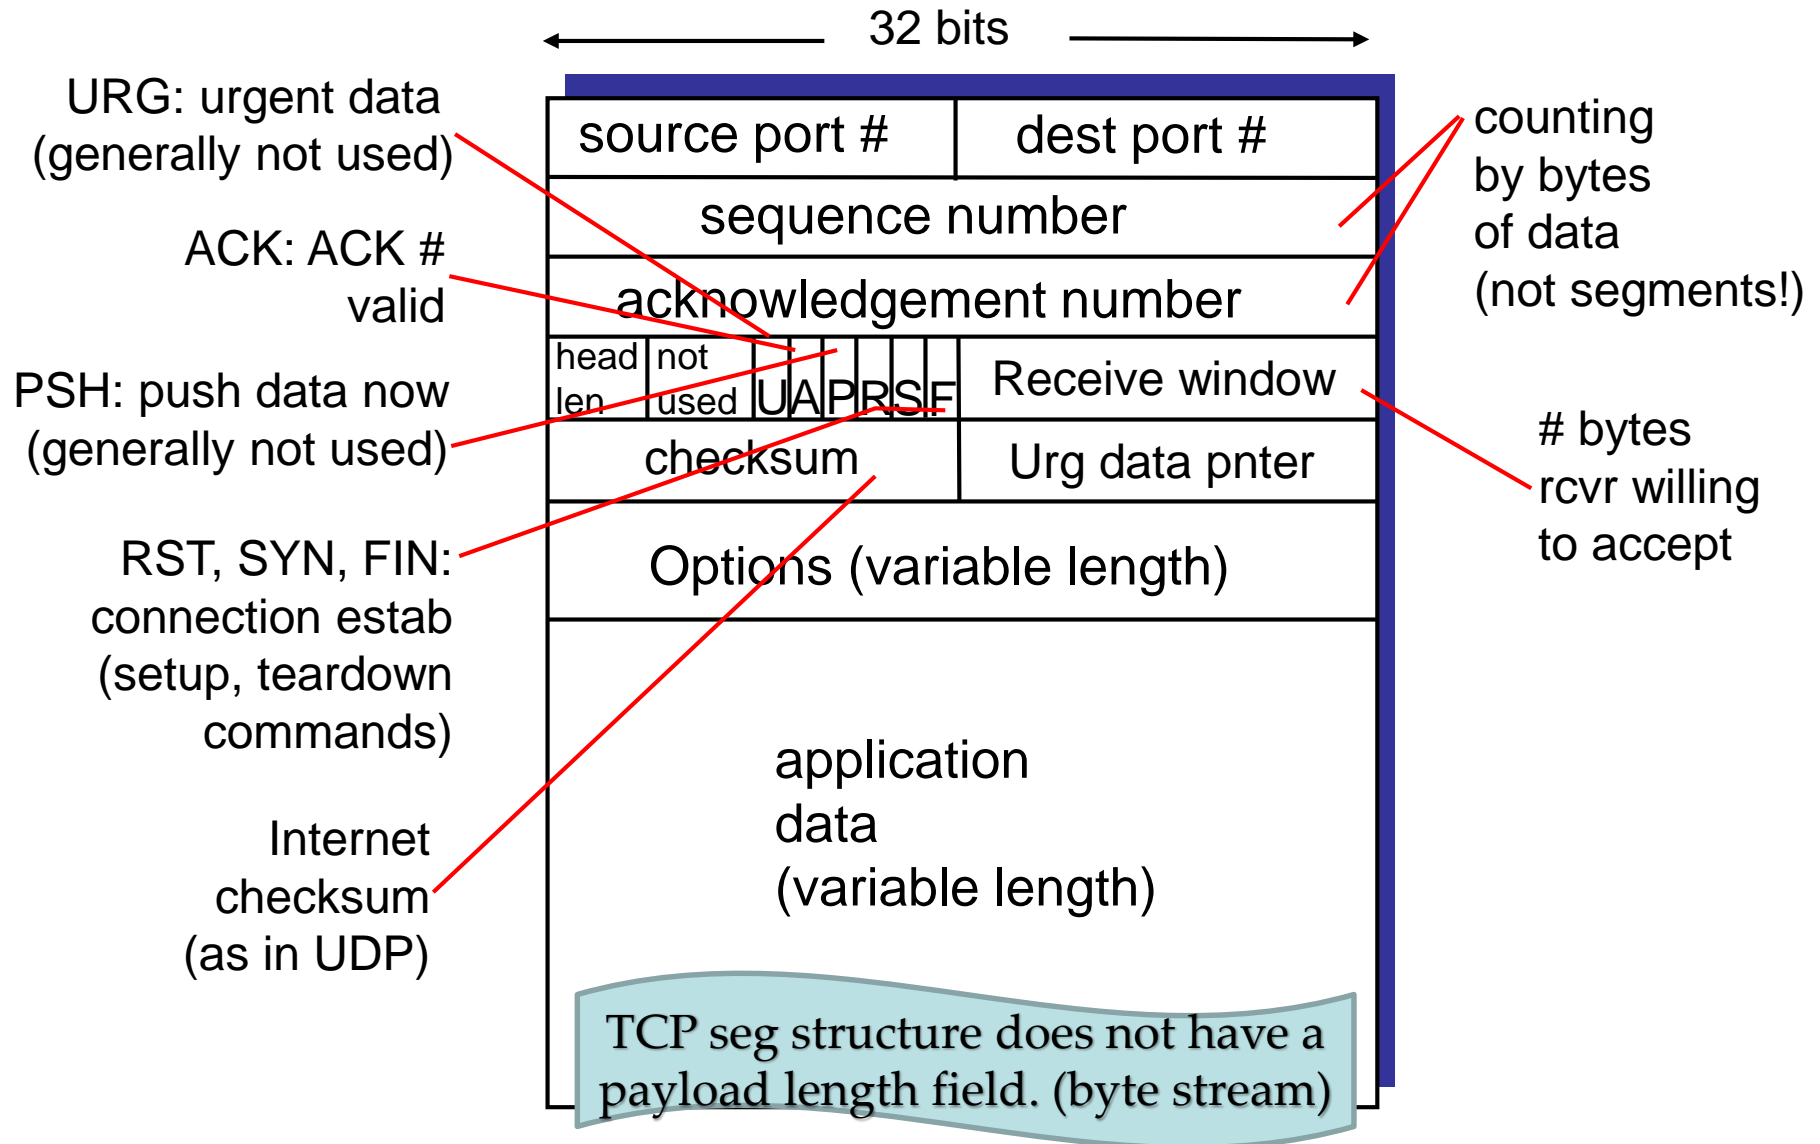
RFCs: 793, 1122, 1323, 2018, 2581

- **point-to-point:**
 - one sender, one receiver
- **reliable, in-order *byte stream*:**
 - no (app. layer) “message boundaries”
- **pipelined:**
 - TCP congestion and flow control set window size



- **full duplex data:**
 - bi-directional data flow in same connection
 - MSS: maximum segment size
- **connection-oriented:**
 - handshaking (exchange of control msgs) init's sender, receiver state before data exchange
- **flow controlled:**
 - sender will not overwhelm receiver

TCP segment structure



No.	Time	Source	Destination	Protocol	Length	Info
1	0.000000	10.34.40.169	132.181.107.25	TCP	66	50769→25 [SYN] Seq=3883811340 Win=8192
2	0.001002	132.181.107.25	10.34.40.169	TCP	66	25→50769 [SYN, ACK] Seq=1828797403 Ack=
3	0.001049	10.34.40.169	132.181.107.25	TCP	54	50769→25 [ACK] Seq=3883811341 Ack=18287
4	0.002254	132.181.107.25	10.34.40.169	SMTP	161	S: 220 UCXHUBCAS01-D.canterbury.ac.nz
5	0.199515	10.34.40.169	132.181.107.25	TCP	54	50769→25 [ACK] Seq=3883811341 Ack=18287
6	300.488449	132.181.107.25	10.34.40.169	SMTP	98	S: 451 4.7.0 Timeout waiting for client
7	300.488450	132.181.107.25	10.34.40.169	TCP	54	25→50769 [FIN, ACK] Seq=1828797555 Ack=
8	300.488515	10.34.40.169	132.181.107.25	TCP	54	50769→25 [ACK] Seq=3883811341 Ack=18287
9	300.489262	10.34.40.169	132.181.107.25	TCP	54	50769→25 [FIN, ACK] Seq=3883811341 Ack=
10	300.490656	132.181.107.25	10.34.40.169	TCP	54	25→50769 [ACK] Seq=1828797556 Ack=38838

▷ Frame 2: 66 bytes on wire (528 bits), 66 bytes captured (528 bits) on interface 0

▷ Ethernet II, Src: JuniperN_ef:61:00 (2c:21:31:ef:61:00), Dst: IntelCor_b6:fe:63 (80:19:34:b6:fe:63)

▷ Internet Protocol Version 4, Src: 132.181.107.25, Dst: 10.34.40.169

▣ Transmission Control Protocol, Src Port: 25, Dst Port: 50769, Seq: 1828797403, Ack: 3883811341, Len: 0

Source Port: 25

Destination Port: 50769

[Stream index: 0]

[TCP Segment Len: 0]

Sequence number: 1828797403

Acknowledgment number: 3883811341

Header Length: 32 bytes

▷ Flags: 0x012 (SYN, ACK)

Window size value: 8192

[Calculated window size: 8192]

Checksum: 0x9393 [unverified]

[Checksum Status: Unverified]

Urgent pointer: 0

▣ Options: (12 bytes), Maximum segment size, No-Operation (NOP), Window scale, No-Operation (NOP), No-Operation (NOP), SACK permitted

▷ Maximum segment size: 1460 bytes

▷ No-Operation (NOP)

▷ Window scale: 8 (multiply by 256)

▷ No-Operation (NOP)

▷ No-Operation (NOP)

▷ TCP SACK Permitted Option: True

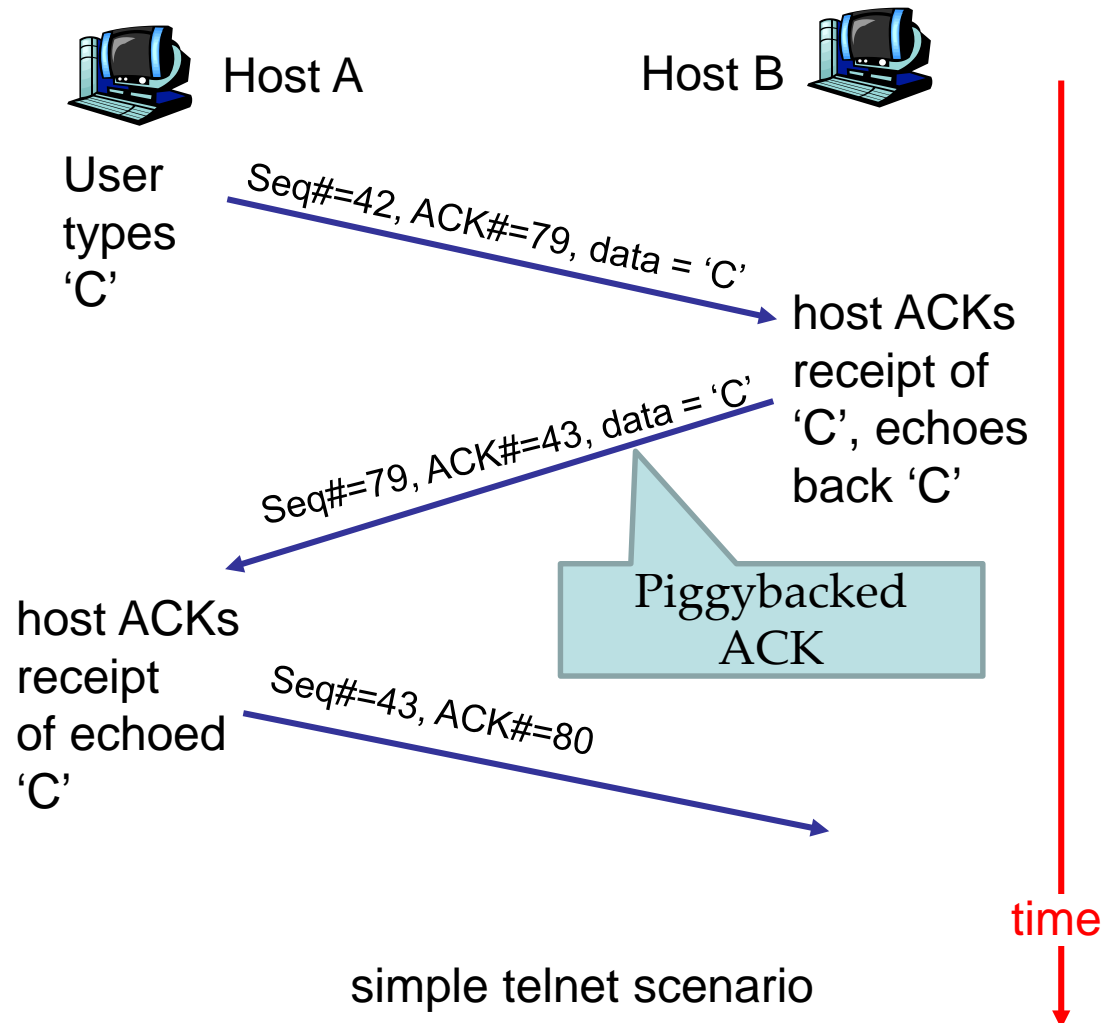
TCP seq. #'s and ACKs

Seq. #'s:

- byte stream
“number” of first
byte in segment's
data

ACKs:

- seq # of next byte
expected from other
side
- cumulative ACK



TCP Connection Management

Three way handshake:

Step 1: client host sends TCP SYN segment (*SYN bit is set to 1*) to server

- specifies initial seq # (*randomly chosen*), client_isn;
- no data

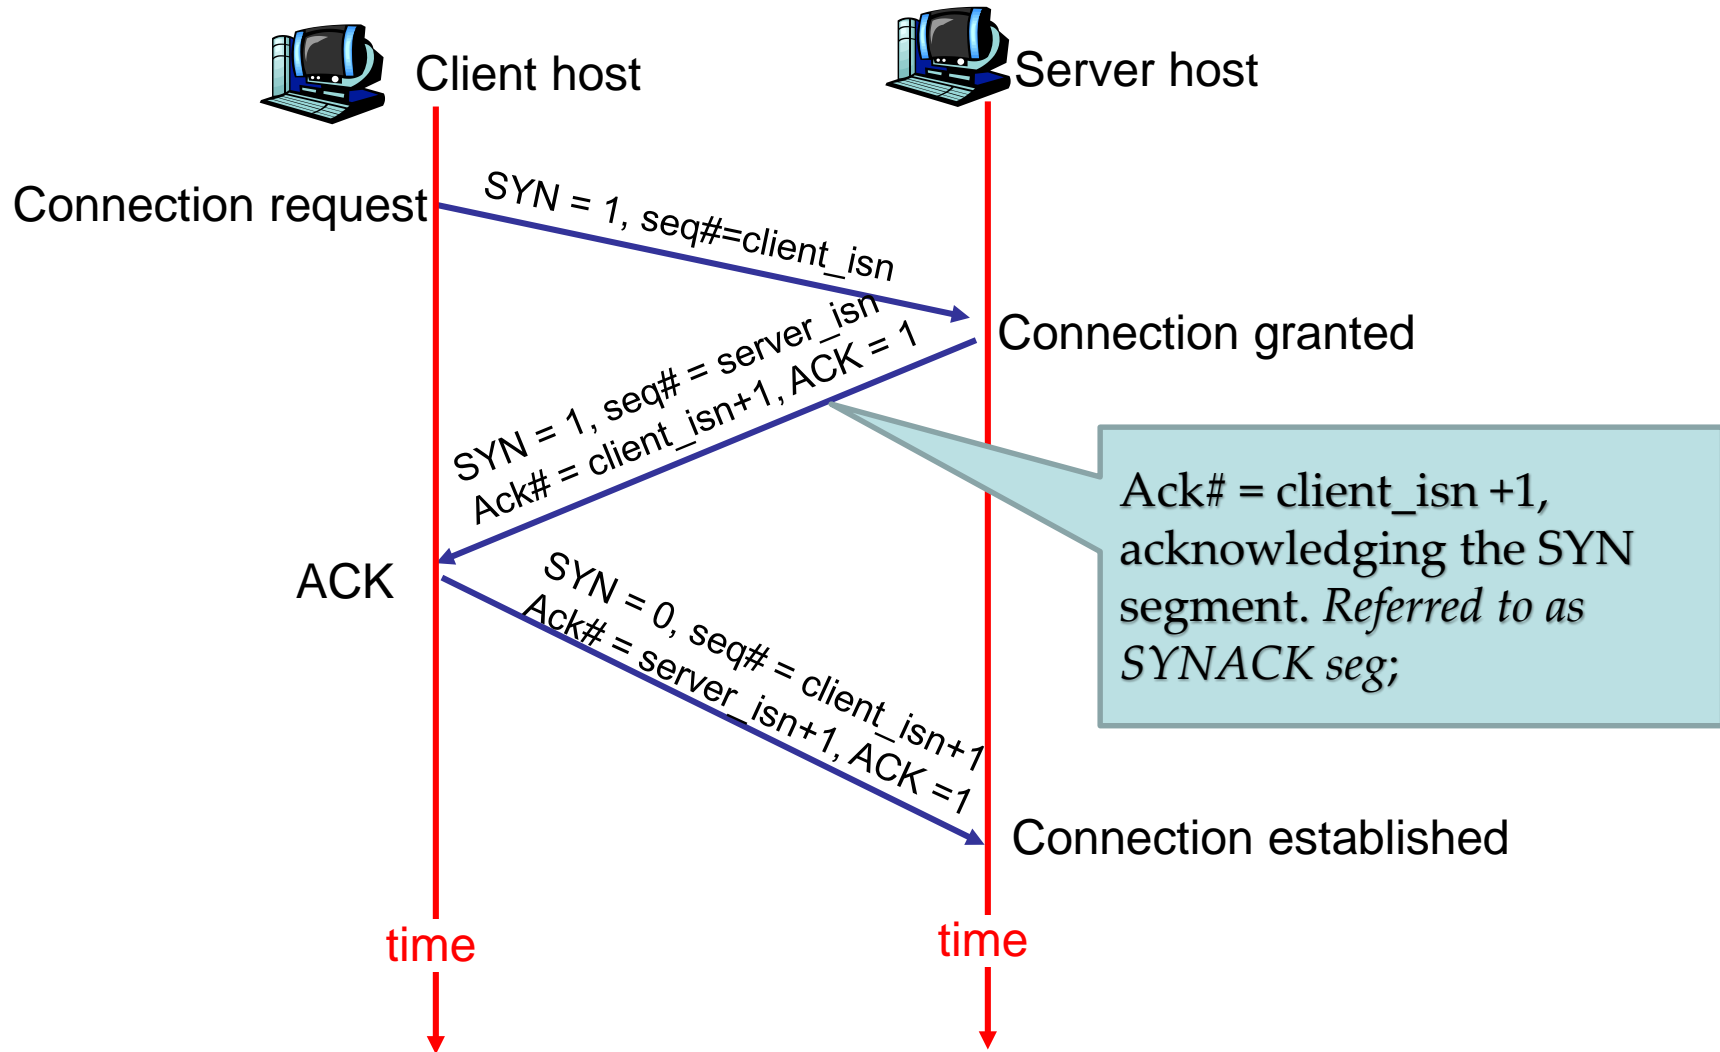
Step 2: server host receives SYN, replies with SYNACK segment

- server allocates buffers
- specifies server initial seq. #, server_isn;
- SYN = 1; ACK# = client_isn + 1; ACK = 1; seq# = server_isn;

Step 3: client receives SYNACK, replies with ACK segment, which may contain data;

SYN = 0; ACK = 1; ACK# = server_isn + 1; seq# = client_isn + 1;

TCP 3-way handshaking



No.	Time	Source	Destination	Protocol
1	0.000000	192.168.88.161	128.238.26.21	TCP
2	0.217170	128.238.26.21	192.168.88.161	TCP
3	0.217326	192.168.88.161	128.238.26.21	TCP
4	51.906622	128.238.26.21	192.168.88.161	TCP
5	51.907091	192.168.88.161	128.238.26.21	TCP
6	51.907277	128.238.26.21	192.168.88.161	TCP

Frame 1: 74 bytes on wire (592 bits), 74 bytes captured (592 bits) on interface 0

Ethernet II, Src: Vmware_b9:06:32 (00:0c:29:b9:06:32), Dst: Vmware_e5:ea:29 (00:0c:29:b9:06:29)

Internet Protocol Version 4, Src: 192.168.88.161, Dst: 128.238.26.21

Transmission Control Protocol, Src Port: 44058, Dst Port: 80, Seq: 2649682993, Len: 0

Source Port: 44058
Destination Port: 80
[Stream index: 0]
[TCP Segment Len: 0]
Sequence number: 2649682993
Acknowledgment number: 0
Header Length: 40 bytes
Flags: 0x0000 (SYN)
Window size value: 29200
[Calculated window size: 29200]
Checksum: 0xcb45 [unverified]
[Checksum Status: Unverified]
Urgent pointer: 0
Options: (20 bytes), Maximum segment size 65535

client_isn

Frame 2: 58 bytes on wire (464 bits), 58 bytes captured (464 bits) on interface 0

Ethernet II, Src: Vmware_e5:ea:29 (00:0c:29:b9:06:29), Dst: Vmware_b9:06:32 (00:0c:29:b9:06:32)

Internet Protocol Version 4, Src: 128.238.26.21, Dst: 192.168.88.161

Transmission Control Protocol, Src Port: 80, Dst Port: 44058, Seq: 2058359514, Len: 0

Source Port: 80
Destination Port: 44058
[Stream index: 0]
[TCP Segment Len: 0]
Sequence number: 2058359514
Acknowledgment number: 2649682994
Header Length: 24 bytes
Flags: 0x0100 (SYN, ACK)
Window size value: 64240
[Calculated window size: 64240]
Checksum: 0x1cc2 [unverified]
[Checksum Status: Unverified]
Urgent pointer: 0
Options: (4 bytes), Maximum segment size 65535
[SEQ/ACK analysis]

server_isn

client_isn + 1

Frame 3: 54 bytes on wire (432 bits), 54 bytes captured (432 bits) on interface 0

Ethernet II, Src: Vmware_b9:06:32 (00:0c:29:b9:06:32), Dst: Vmware_e5:ea:29 (00:0c:29:b9:06:29)

Internet Protocol Version 4, Src: 192.168.88.161, Dst: 128.238.26.21

Transmission Control Protocol, Src Port: 44058, Dst Port: 80, Seq: 2649682994, Len: 0

Source Port: 44058
Destination Port: 80
[Stream index: 0]
[TCP Segment Len: 0]
Sequence number: 2649682994
Acknowledgment number: 2058359515
Header Length: 20 bytes
Flags: 0x0100 (ACK)
Window size value: 29200
[Calculated window size: 29200]
[Window size scaling factor: -2 (no window scaling used)]
Checksum: 0xbd5f [unverified]
[Checksum Status: Unverified]
Urgent pointer: 0
Options: (0 bytes)
[SEQ/ACK analysis]

client_isn + 1

server_isn + 1

TCP Connection Management (closing)


All good things must come to an end:

Step 1: client end system sends TCP FIN control segment to server

Step 2: server receives FIN, replies with ACK. Closes connection, sends FIN.

Step 3: client receives FIN, replies with ACK.

Step 4: server, receives ACK. Connection closed.



37	160.494102	10.34.40.169	132.181.107.25	TCP	54	57483→25	[FIN, ACK]	Seq=3400705033	Ack=1545042735	Win=6...
38	160.495931	132.181.107.25	10.34.40.169	TCP	54	25→57483	[ACK]	Seq=1545042735	Ack=3400705034	Win=131328...
39	160.496329	132.181.107.25	10.34.40.169	TCP	54	25→57483	[FIN, ACK]	Seq=1545042735	Ack=3400705034	Win=1...
40	160.496413	10.34.40.169	132.181.107.25	TCP	54	57483→25	[ACK]	Seq=3400705034	Ack=1545042736	Win=65536 ...

Internet Protocol Version 4, Src: 10.34.40.169, Dst: 132.181.107.25

Transmission Control Protocol, Src Port: 57483, Dst Port: 25, Seq: 3400705033, Ack: 1545042735, Len: 0

Source Port: 57483

Destination Port: 25

[Stream index: 1]

[TCP Segment Len: 0]

Sequence number: 3400705033

Acknowledgment number: 1545042735

Header Length: 20 bytes

Flags: 0x011 (FIN, ACK)

000. = Reserved: Not set
 ...0 = Nonce: Not set
0... = Congestion Window Reduced
0... = ECN-Echo: Not set
0... = Urgent: Not set
1... = Acknowledgment: Set
0... = Push: Not set
0... = Reset: Not set
0... = Syn: Not set

...1 = Fin: Set

[TCP Flags:A...F]

Window size value: 256

[Calculated window size: 65536]

[Window size scaling factor: 256]

Checksum: 0x6592 [unverified]

[Checksum Status: Unverified]

Urgent pointer: 0

Frame 38: 54 bytes on wire (432 bits), 54 bytes captured (432 bits) on interface 0

Ethernet II, Src: JuniperN_ef:61:00 (2c:21:31:ef:61:00), Dst: IntelCor_b6:fe:63 (80:19:34:b6:fe:63)

Internet Protocol Version 4, Src: 132.181.107.25, Dst: 10.34.40.169

Transmission Control Protocol, Src Port: 25, Dst Port: 57483, Seq: 1545042735, Ack: 3400705034, Len: 0

Source Port: 25

Destination Port: 57483

[Stream index: 1]

[TCP Segment Len: 0]

Sequence number: 1545042735

Acknowledgment number: 3400705034

Header Length: 20 bytes

Flags: 0x010 (ACK)

000. = Reserved: Not set
 ...0 = Nonce: Not set
0... = Congestion Window Reduced (CWR): Not set
0... = ECN-Echo: Not set
0... = Urgent: Not set
1... = Acknowledgment: Set
0... = Push: Not set
0... = Reset: Not set
0... = Syn: Not set
0... = Fin: Not set


[TCP Flags:A....]

Window size value: 513

[Calculated window size: 131328]

[Window size scaling factor: 256]

Checksum: 0x6491 [unverified]



FIN

ACK

37	160.494102	10.34.40.169	132.181.107.25	TCP
38	160.495931	132.181.107.25	10.34.40.169	TCP
→ 39	160.496329	132.181.107.25	10.34.40.169	TCP
40	160.496413	10.34.40.169	132.181.107.25	TCP

> Frame 39: 54 bytes on wire (432 bits), 54 bytes captured (432 bits) on interface 0
 > Ethernet II, Src: JuniperN_ef:61:00 (2c:21:31:ef:61:00), Dst: IntelCor_b6:fe:63 (80:19:34:b6:fe:63)
 > Internet Protocol Version 4, Src: 132.181.107.25, Dst: 10.34.40.169
 > Transmission Control Protocol, Src Port: 25, Dst Port: 57483, Seq: 1545042735, Ack: 3400705034, Len: 0

Source Port: 25
 Destination Port: 57483
 [Stream index: 1]
 [TCP Segment Len: 0]
 Sequence number: 1545042735
 Acknowledgment number: 3400705034
 Header Length: 20 bytes
 > Flags: 0x011 (FIN, ACK)
 000. = Reserved: Not set
 ...0 = Nonce: Not set
0... = Congestion Window Reduced (CWR): Not set
0... = ECN-Echo: Not set
0... = Urgent: Not set
1... = Acknowledgment: Set
0... = Push: Not set
0... = Reset: Not set
0... = Syn: Not set
 >1... = Fin: Set
 [TCP Flags:A...F]
 Window size value: 513
 [Calculated window size: 131328]
 [Window size scaling factor: 256]
 Checksum: 0x6490 [unverified]


37	160.494102	10.34.40.169	132.181.107.25	TCP
38	160.495931	132.181.107.25	10.34.40.169	TCP
39	160.496329	132.181.107.25	10.34.40.169	TCP
→ 40	160.496413	10.34.40.169	132.181.107.25	TCP

> Frame 40: 54 bytes on wire (432 bits), 54 bytes captured (432 bits) on interface 0
 > Ethernet II, Src: IntelCor_b6:fe:63 (80:19:34:b6:fe:63), Dst: JuniperN_ef:61:00 (2c:21:31:ef:61:00)
 > Internet Protocol Version 4, Src: 10.34.40.169, Dst: 132.181.107.25
 > Transmission Control Protocol, Src Port: 57483, Dst Port: 25, Seq: 3400705034, Ack: 1545042736, Len: 0

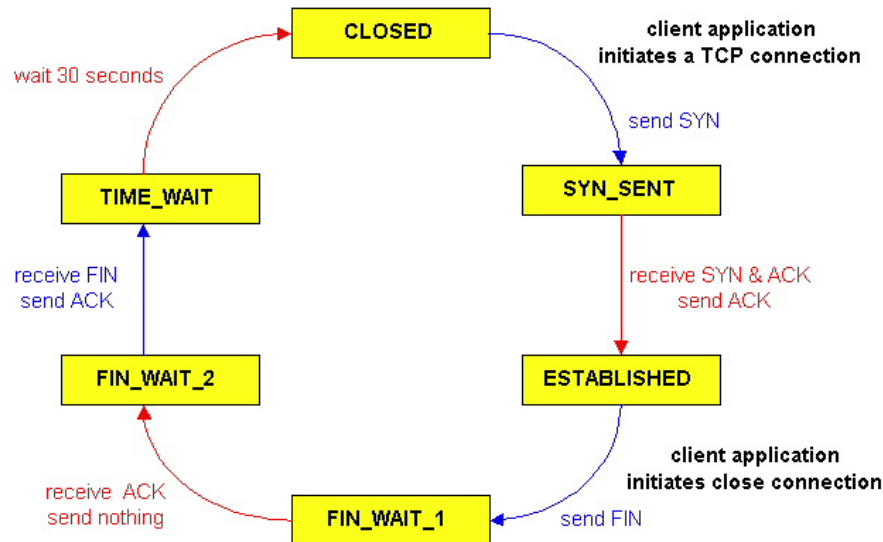
Source Port: 57483
 Destination Port: 25
 [Stream index: 1]
 [TCP Segment Len: 0]
 Sequence number: 3400705034
 Acknowledgment number: 1545042736
 Header Length: 20 bytes
 > Flags: 0x010 (ACK)
 000. = Reserved: Not set
 ...0 = Nonce: Not set
0... = Congestion Window Reduced (CWR): Not set
0... = ECN-Echo: Not set
0... = Urgent: Not set
1... = Acknowledgment: Set
0... = Push: Not set
0... = Reset: Not set
0... = Syn: Not set
0... = Fin: Not set
 [TCP Flags:A....]
 Window size value: 256
 [Calculated window size: 65536]
 [Window size scaling factor: 256]
 Checksum: 0x6591 [unverified]

FIN

ACK



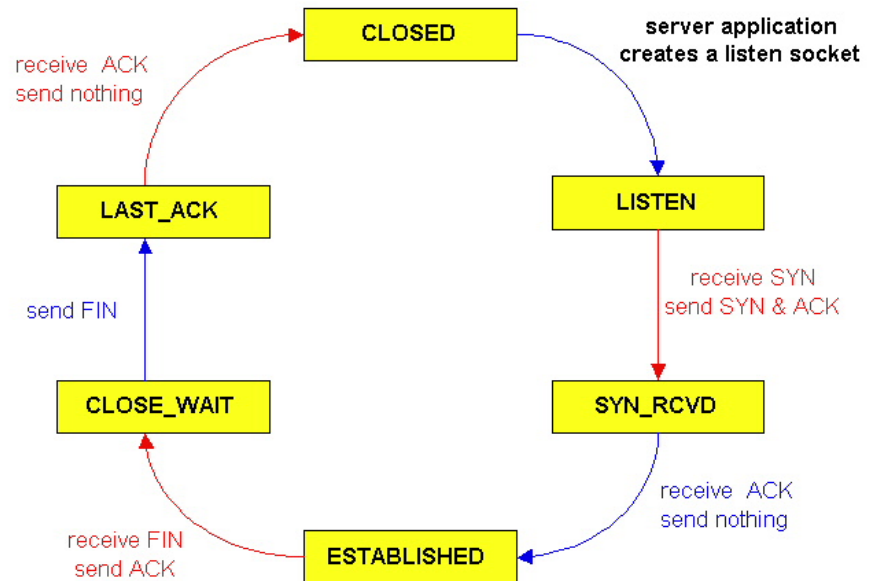
TCP Connection Management (cont)



TCP client lifecycle

This is the normal closing procedure; there are other cases.

TCP server lifecycle



Outline

- Transport-layer services
- Multiplexing and demultiplexing
- Connectionless transport: UDP
- Connection-oriented transport: TCP

References

- [KR3] James F. Kurose, Keith W. Ross, *Computer networking: a top-down approach featuring the Internet*, 3rd edition.
- [PD5] Larry L. Peterson, Bruce S. Davie, *Computer networks: a systems approach*, 5th edition
- [TW5] Andrew S. Tanenbaum, David J. Wetherall, *Computer network*, 5th edition
- [LHBi]Y-D. Lin, R-H. Hwang, F. Baker, *Computer network: an open source approach*, International edition

Acknowledgements

- All slides are developed based on slides from the following two sources:
 - Dr DongSeong Kim's slides for COSC264, University of Canterbury;
 - Prof Aleksandar Kuzmanovic's lecture notes for CS340, Northwestern University,
https://users.cs.northwestern.edu/~akuzma/classes/CS340-w05/lecture_notes.htm